TRIAD: Creating Synergies Between Memory, Disk and Log in Log Structured Key-Value Stores

USENIX ATC '17, Santa Clara CA

O. Balmau — EPFL

D. Didona — EPFL

R. Guerraoui — EPFL

W. Zwaenepoel — EPFL

H. Yuan — Nutanix

A. Arora — Nutanix

K. Gupta — Nutanix

P. Konka — Nutanix

KV Stores

Very **simple** data stores.

KV pairs.

Simple operations: **update, read.**

KV Stores

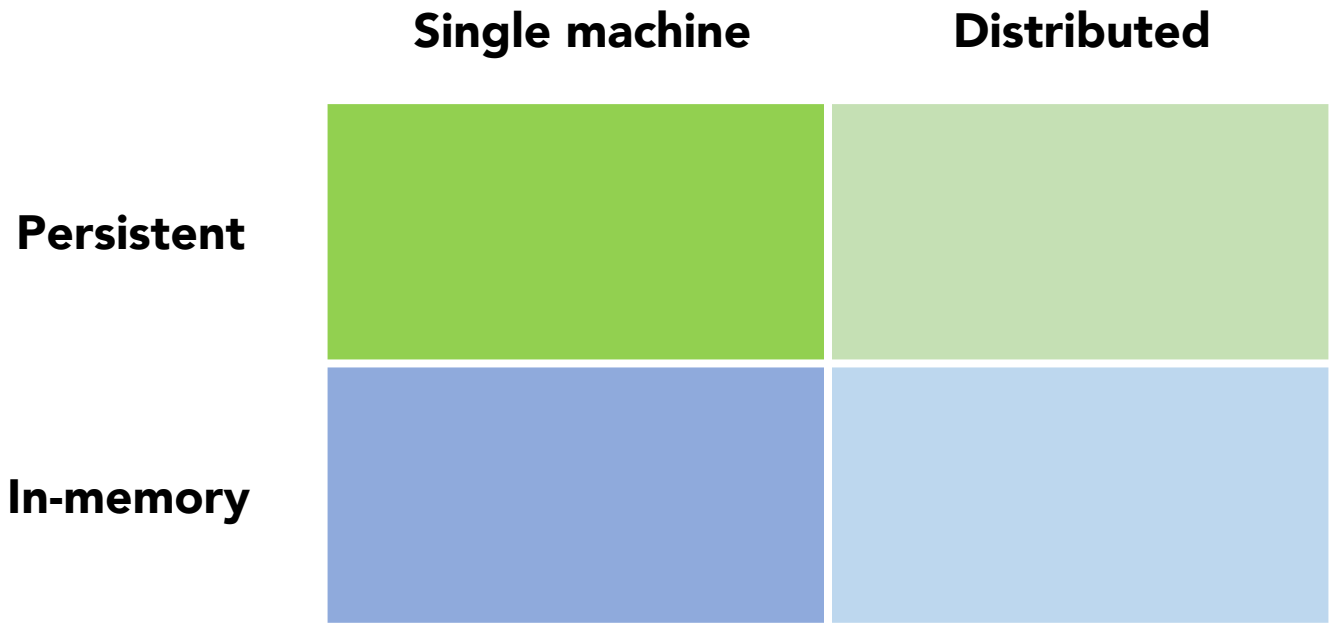

KV Stores

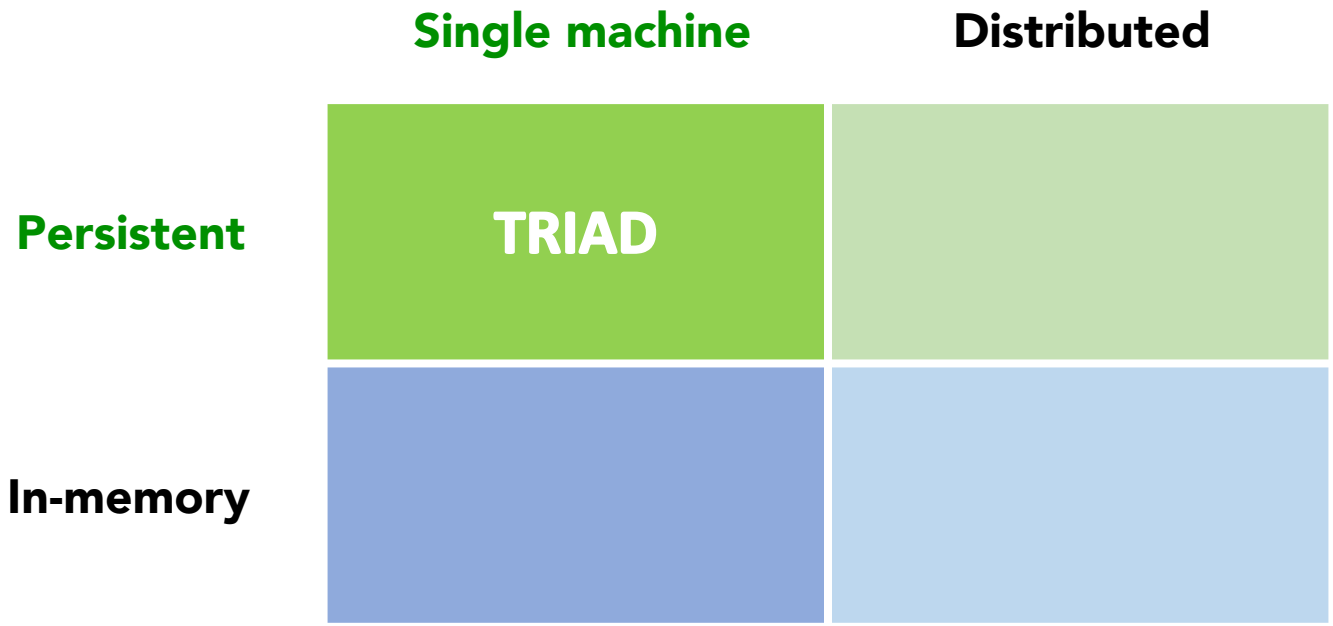

KV Stores

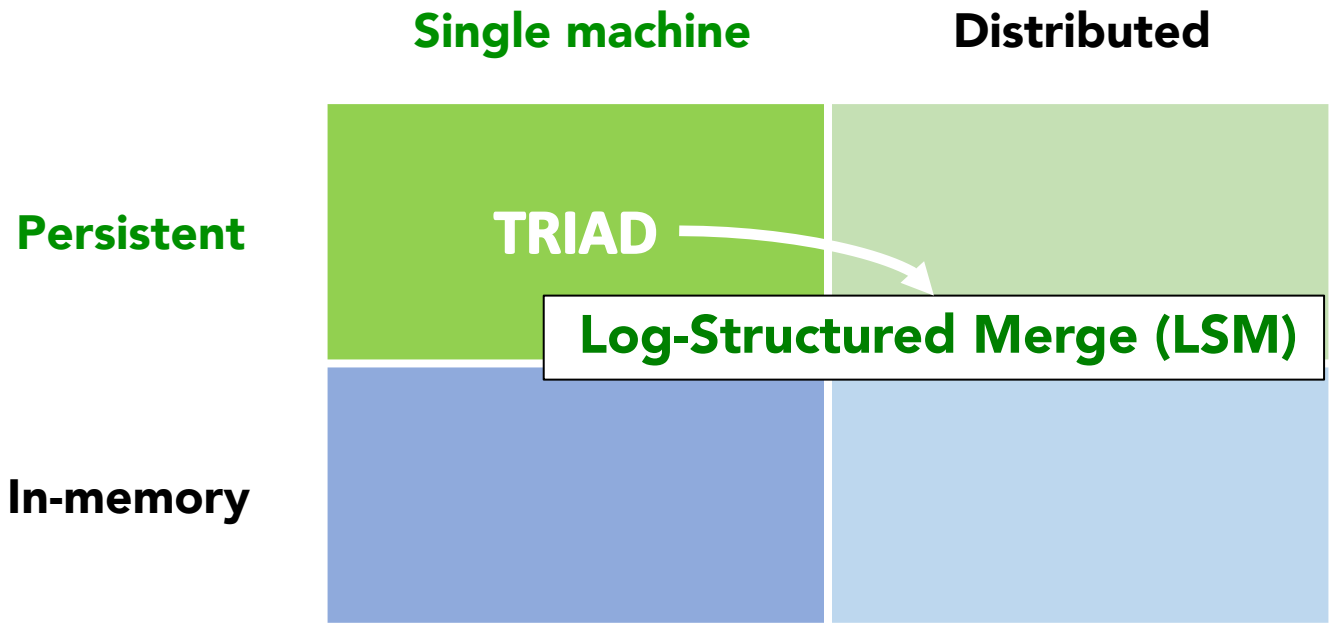
Single machine
Distributed
Persistent
In-memory
TRIAD
Log-Structured Merge (LSM)

TRIAD in a Nutshell

TRIAD LSM KV: achieves 2x throughput on production wklds.

Methods: Reducing background I/O in LSMs.

TRIAD in a Nutshell

TRIAD LSM KV: achieves 2x throughput on production wklds.

Methods: Reducing background I/O in LSMs.

☺ No need to know workload a priori.

☺ LSM KV semantics preserved.

LSM Overview

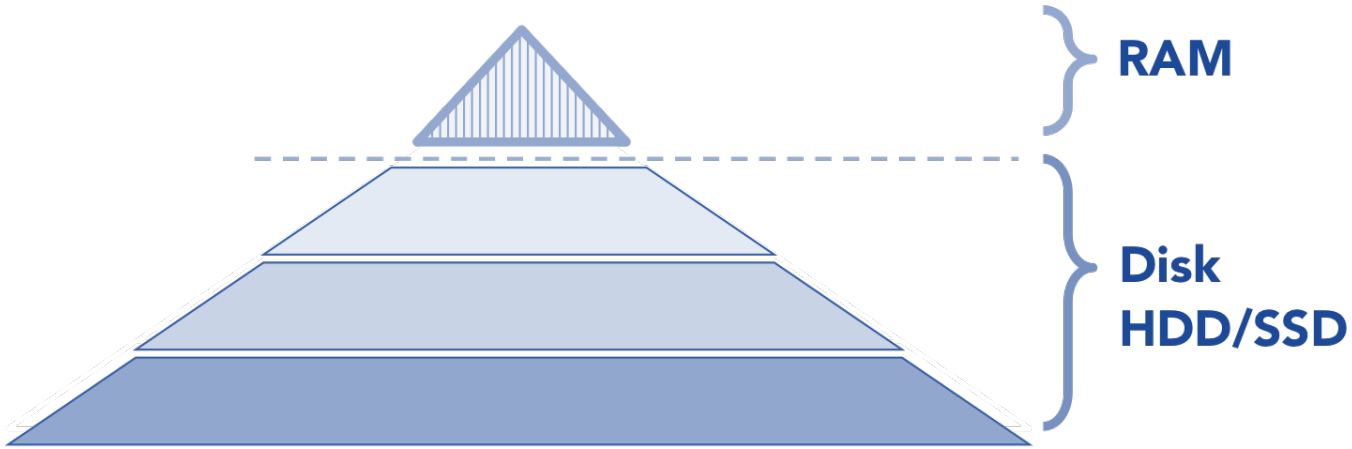

LSM Components

Sorted memory component

SSTables
- sorted files
- many SSTables/Level

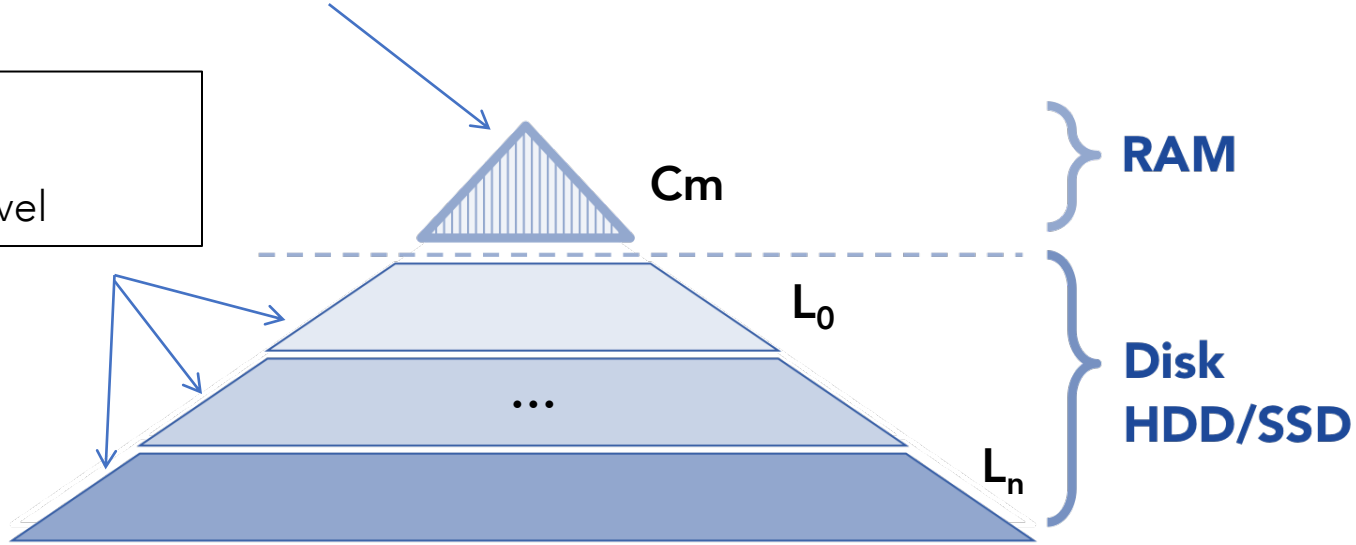

LSM Updates

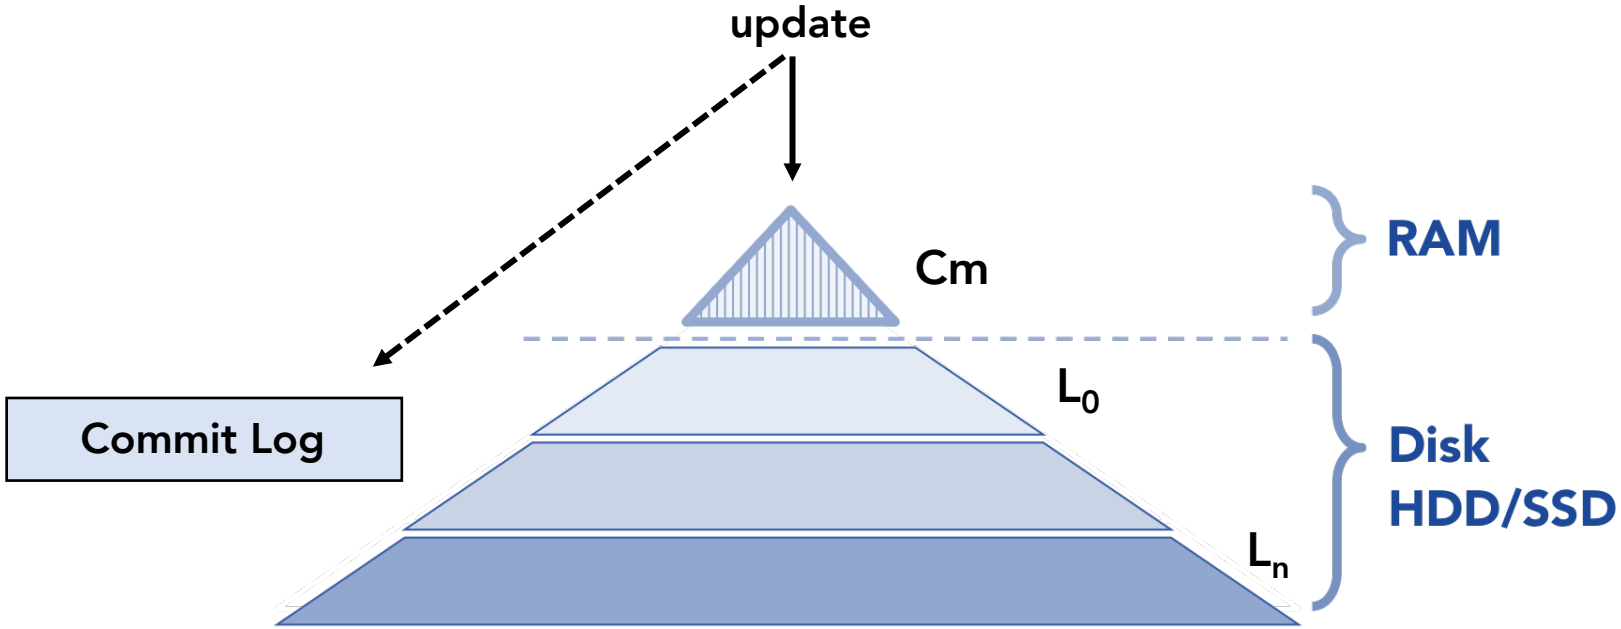

LSM Reads

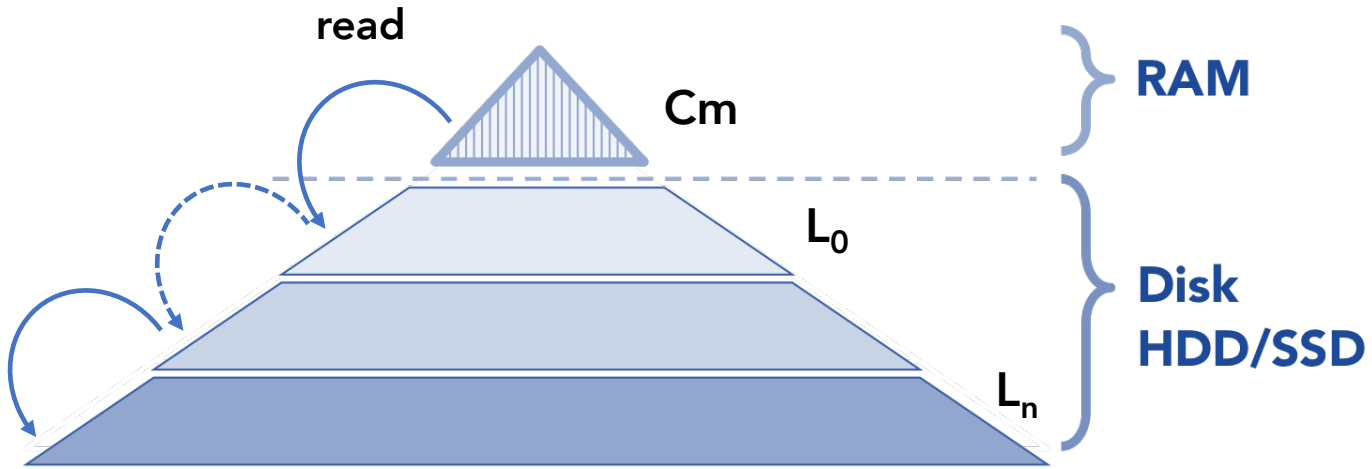

LSM Background Ops: Flushing

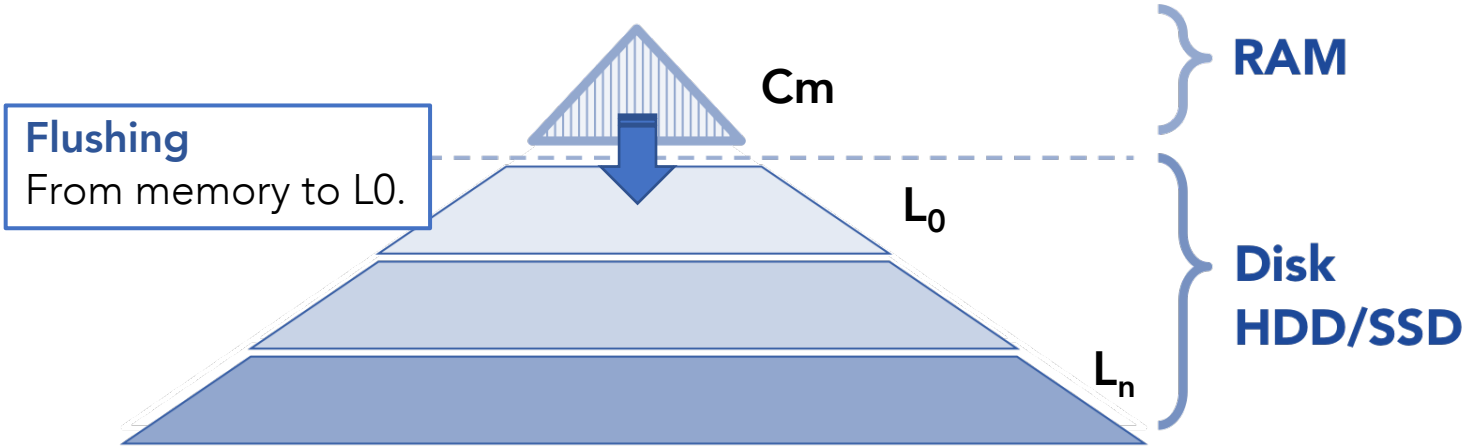

LSM Background Ops: Flushing

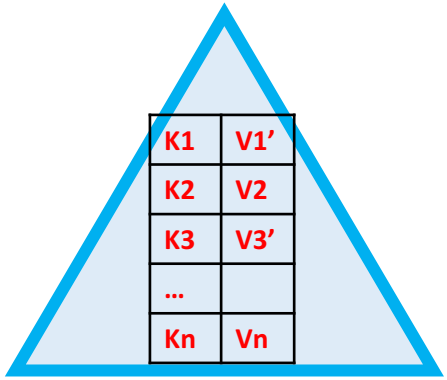

Cm

Mem component full

RAM

Disk

flushing

L0

Commit Log

K1	V1
K2	V2
K1	V1'
K3	V3

LSM Background Ops: Flushing

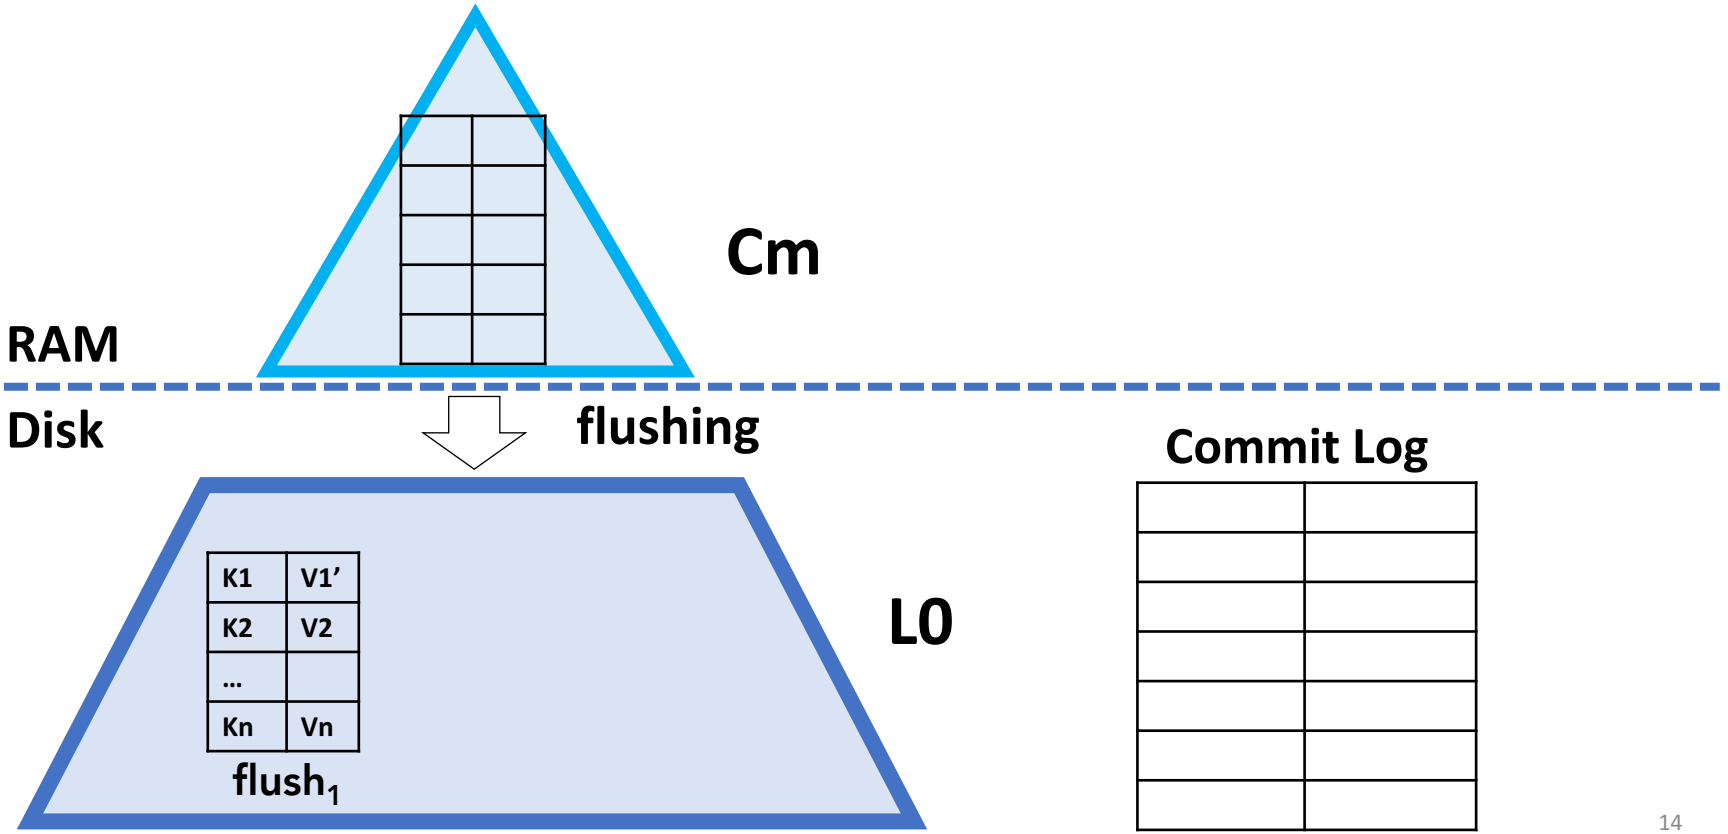

LSM Background Ops: Flushing

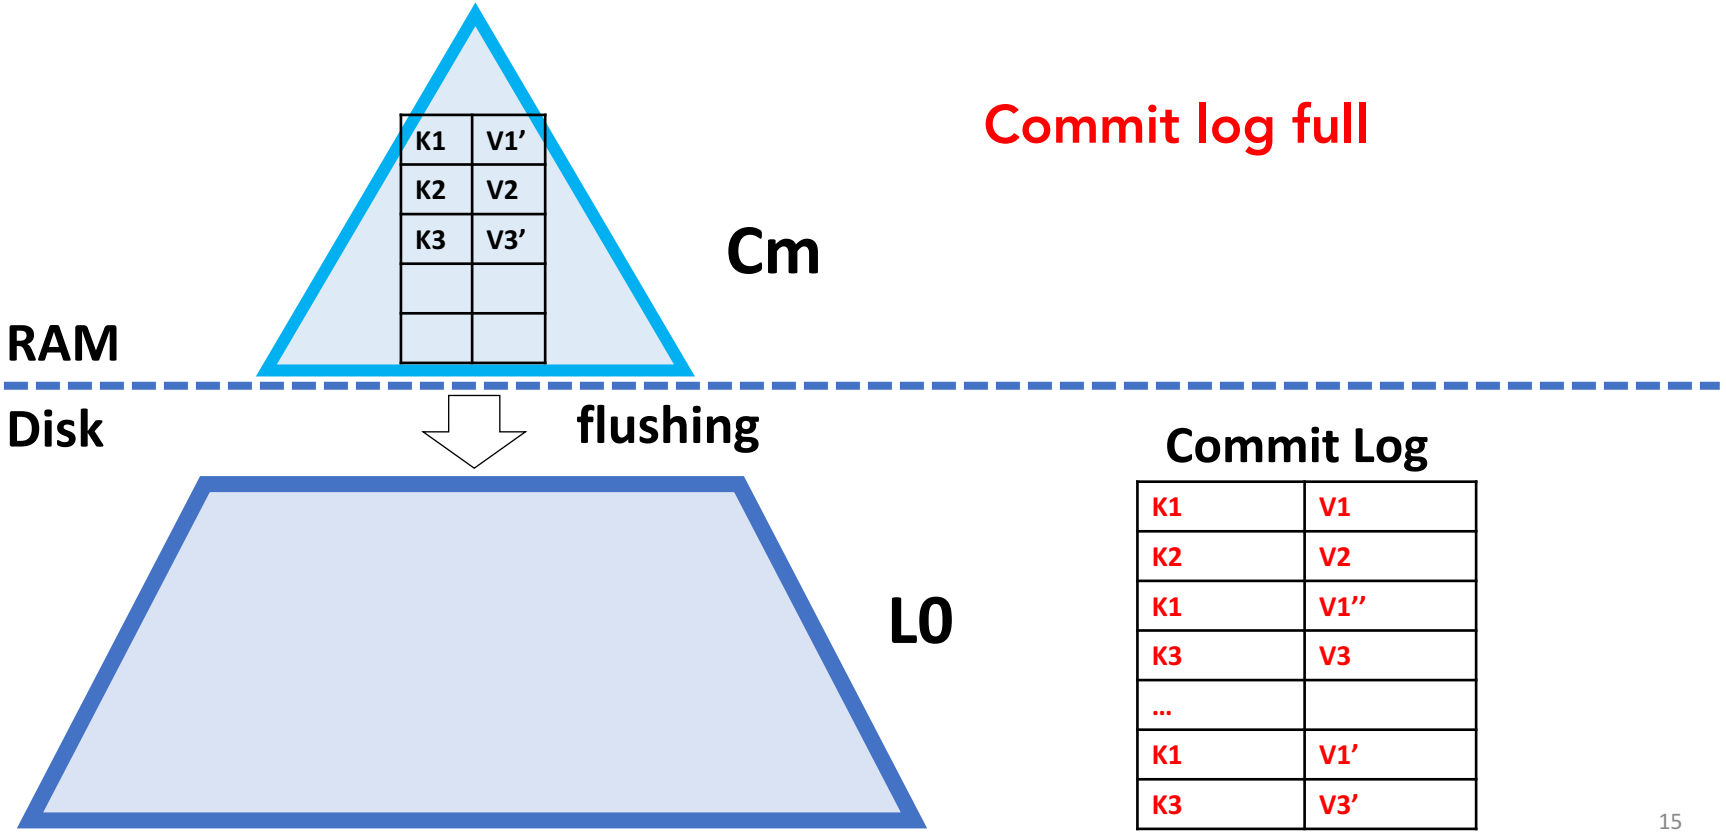

LSM Background Ops: Flushing

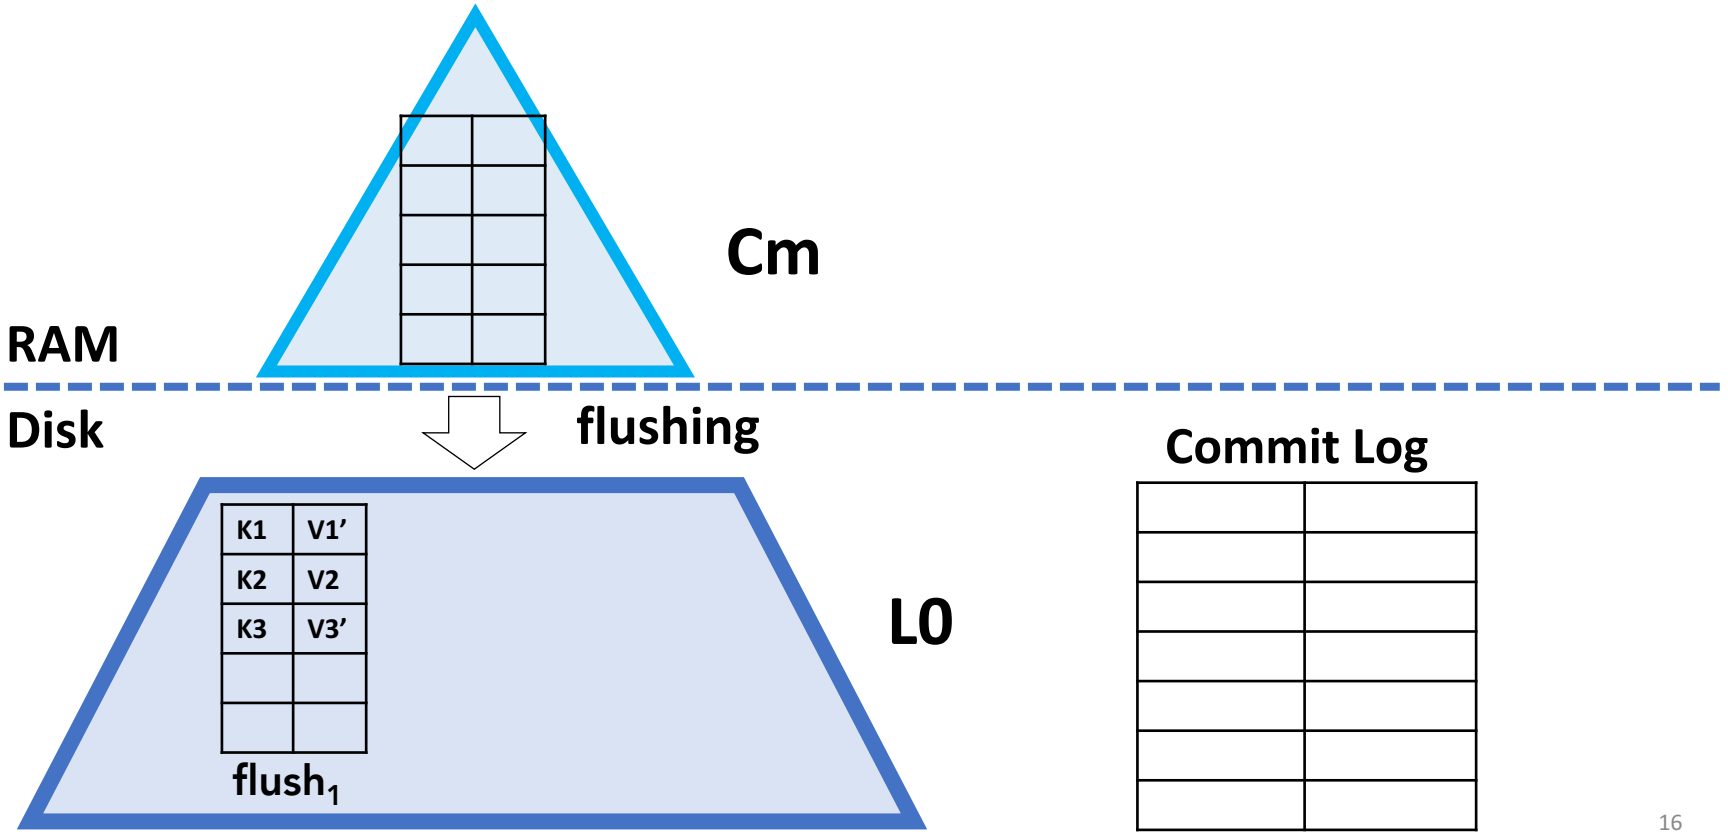

LSM Background Ops: Compaction

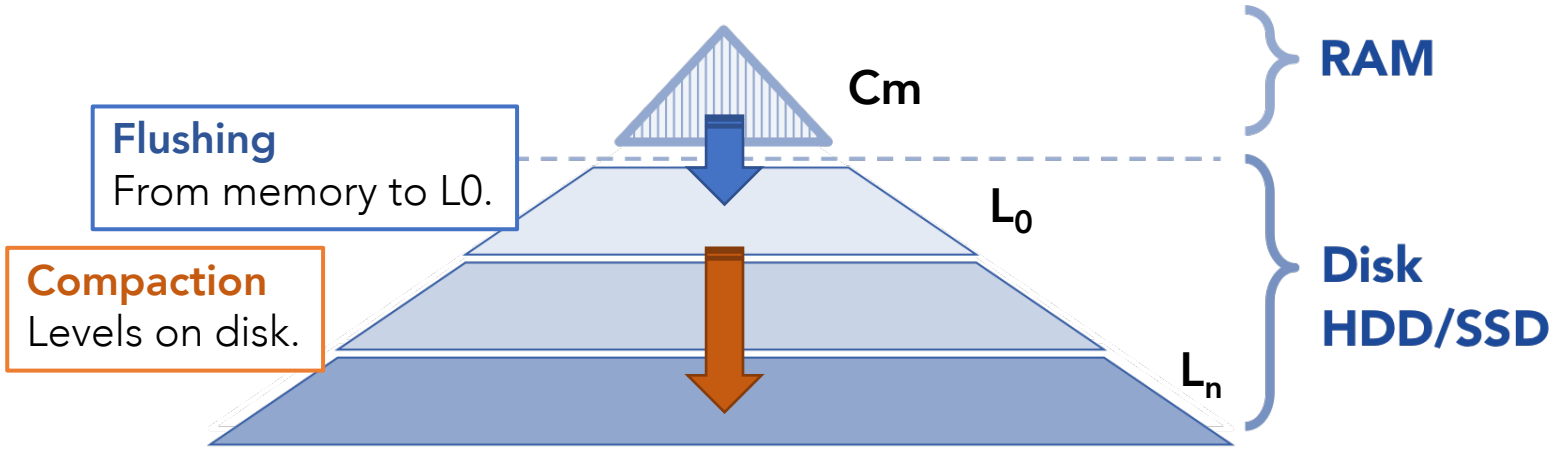

LSM Background Ops: Compaction

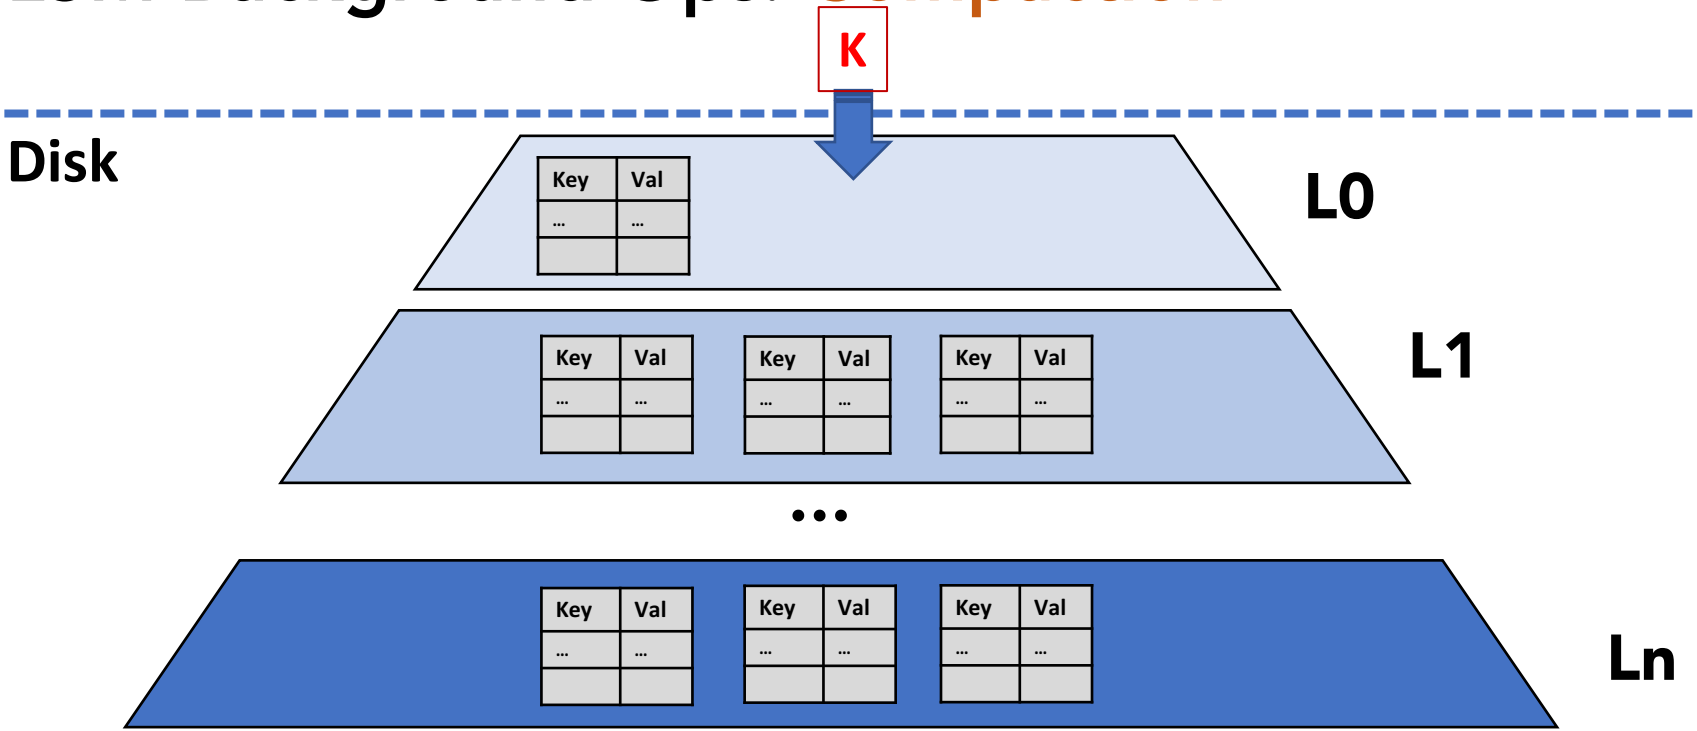

LSM Background Ops: Compaction

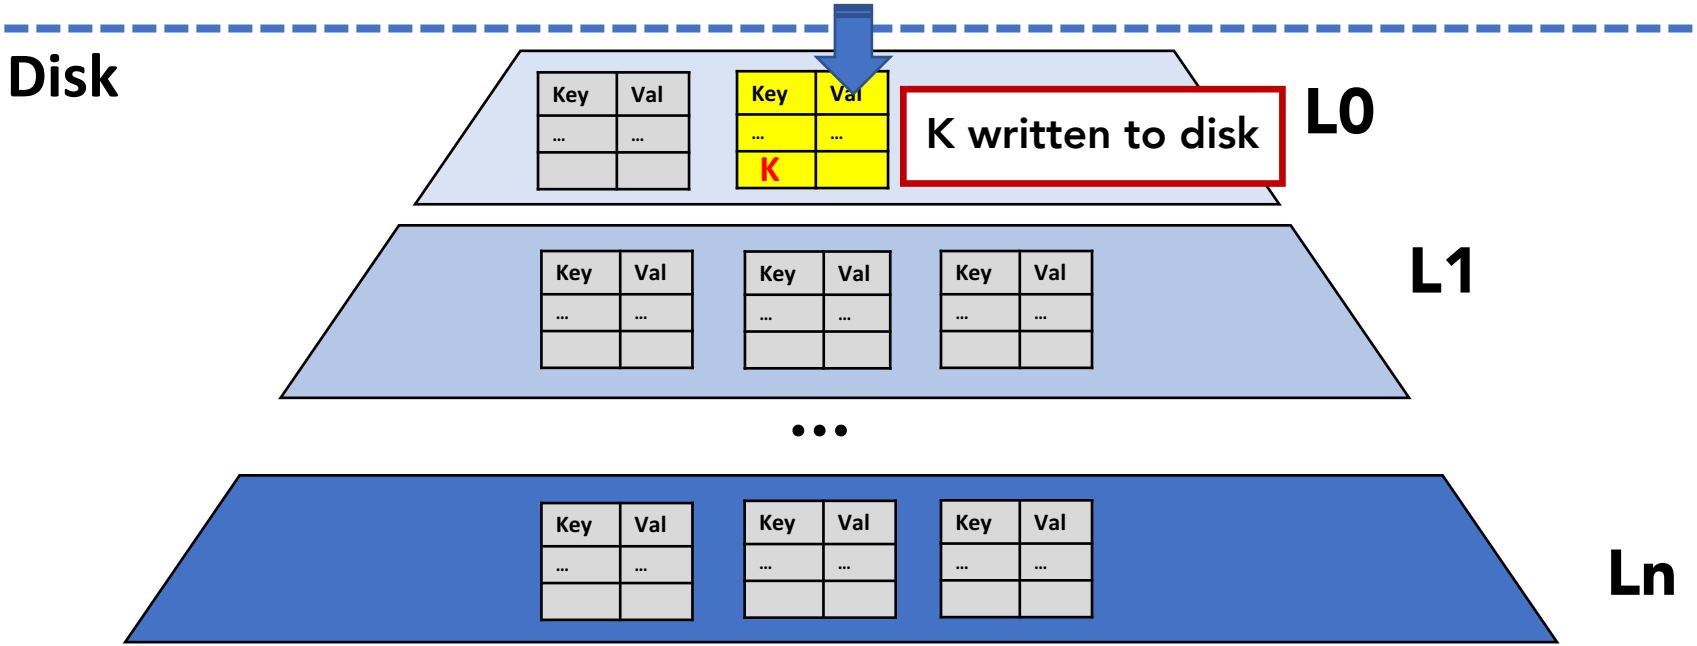

LSM Background Ops: Compaction

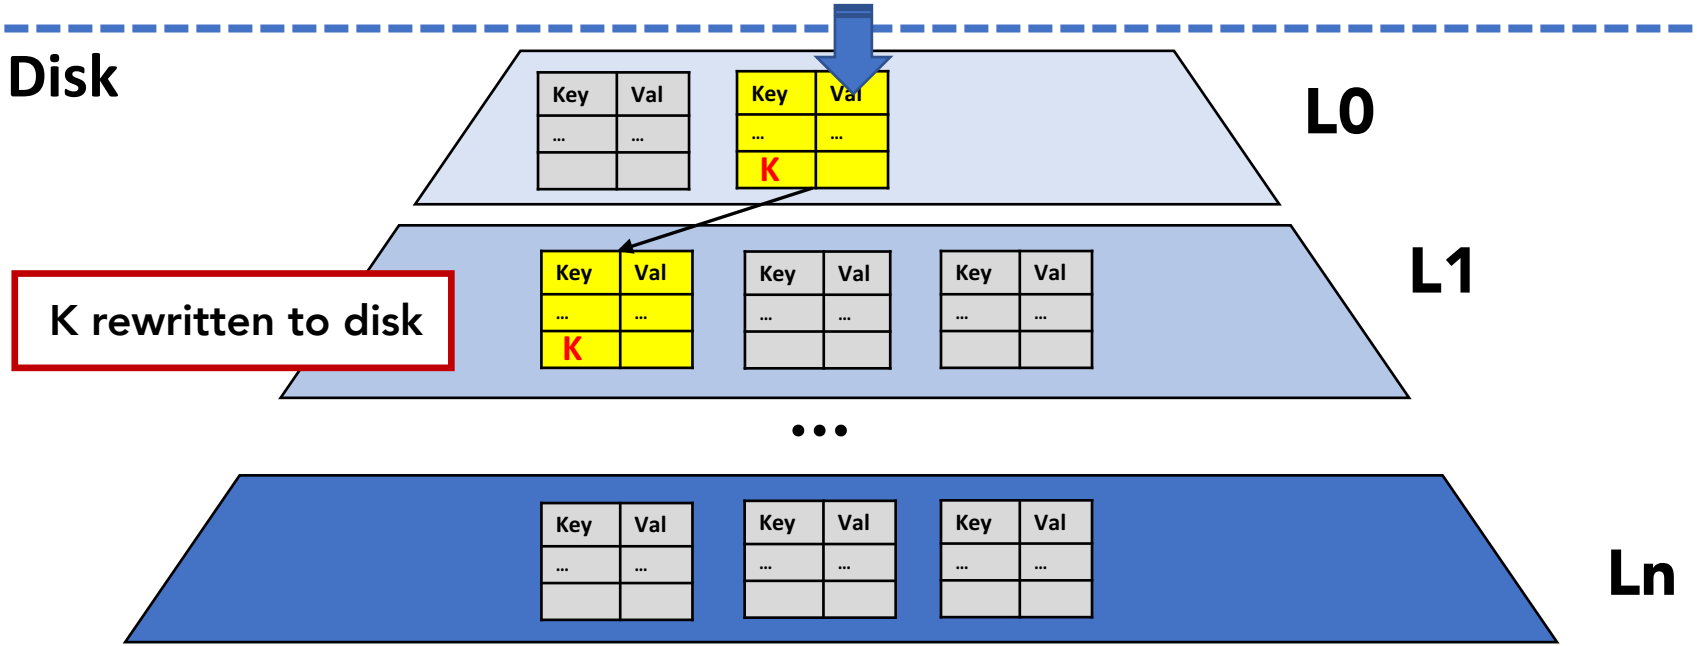

LSM Background Ops: Compaction

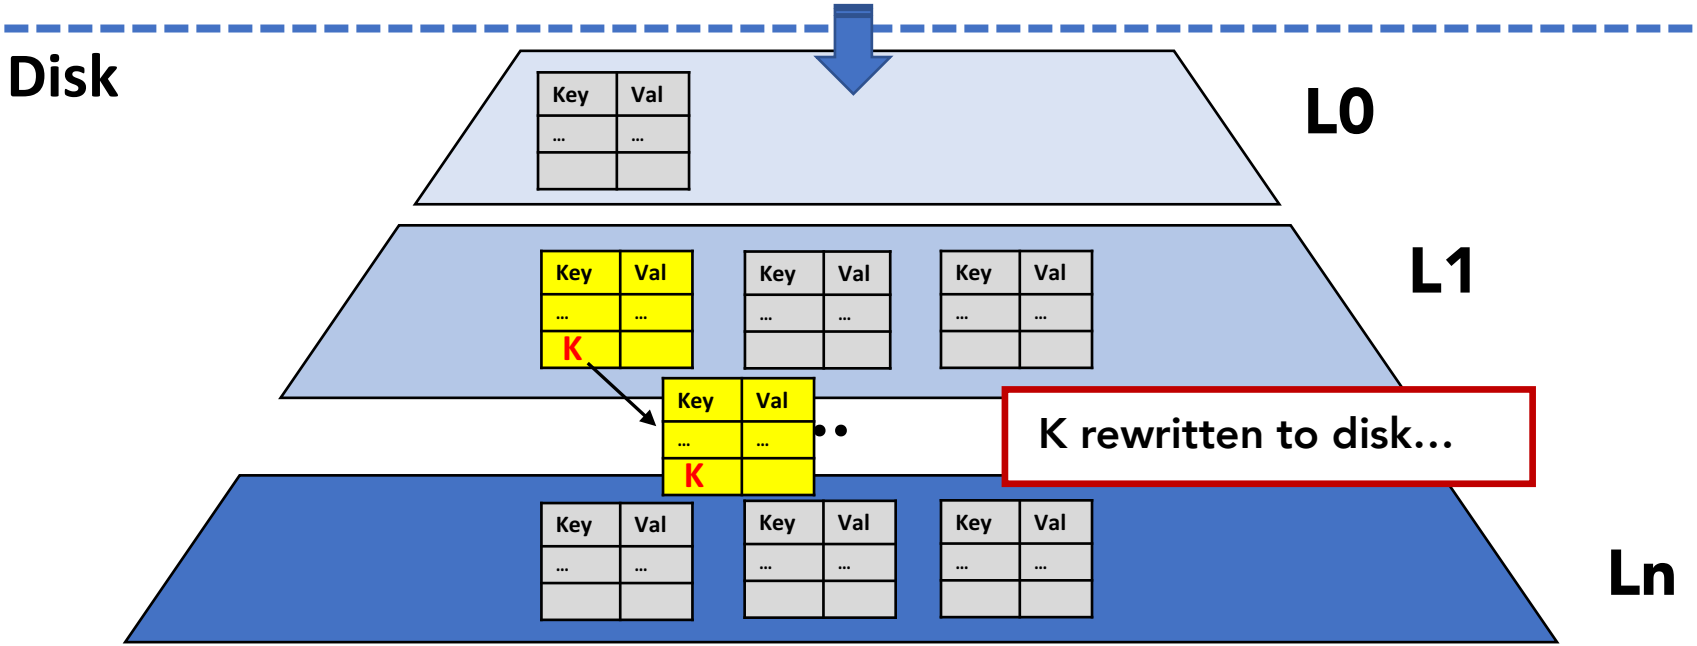

LSM Background Ops: Compaction

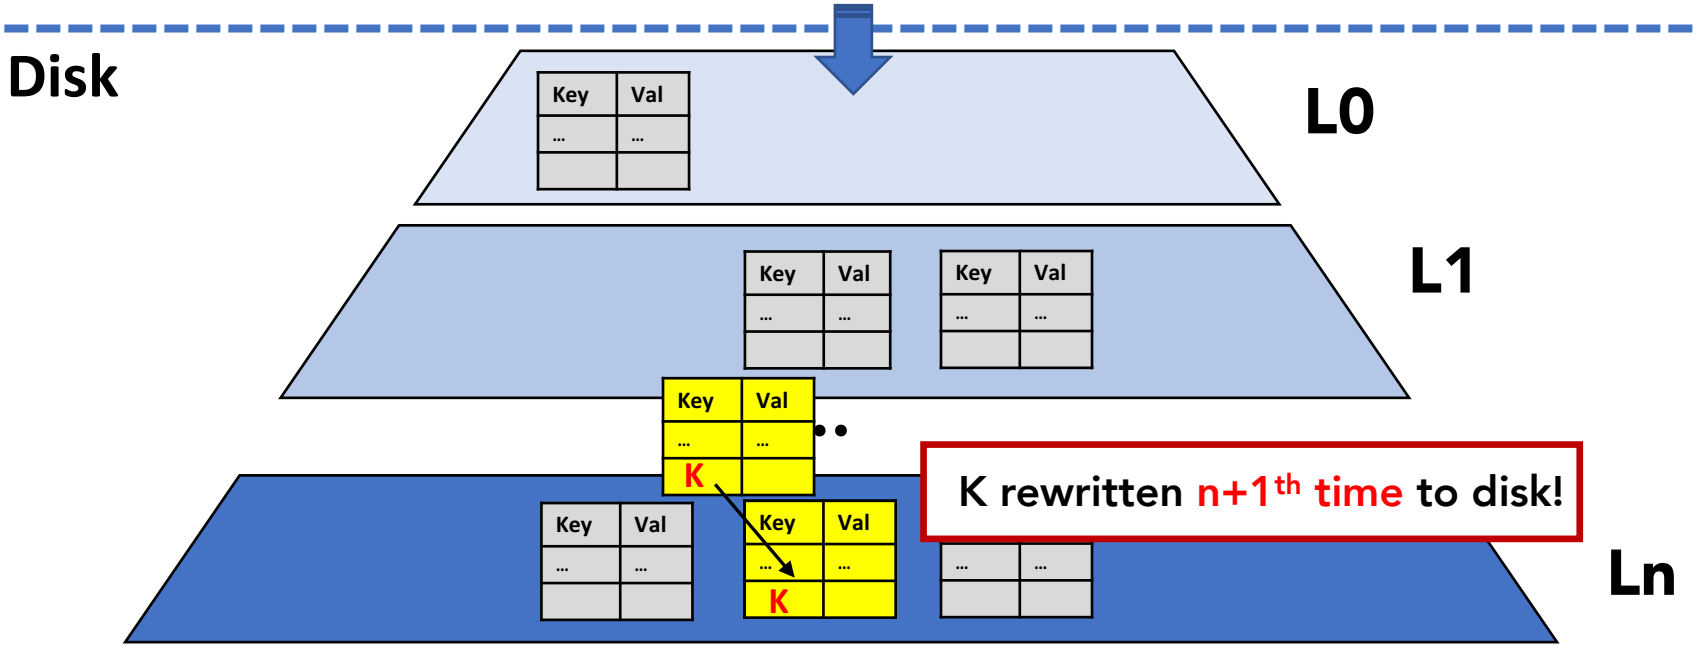

LSM Background Ops: Compaction

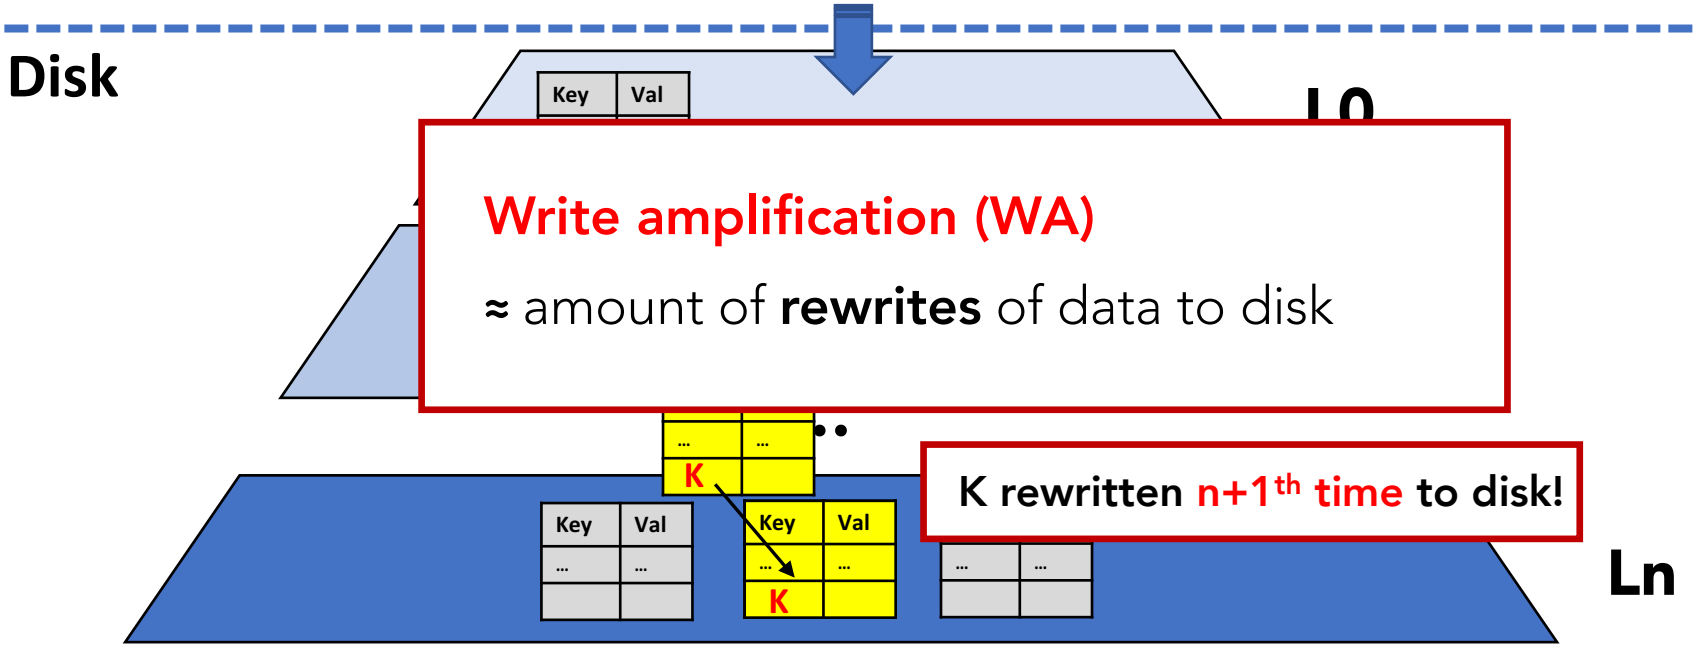

Insight

Severe competition for compute/storage resources between **LSM background ops and user ops.**

Background I/O Overhead

- **Long & slow bg. ops ⟶ slowdown of user ops.**

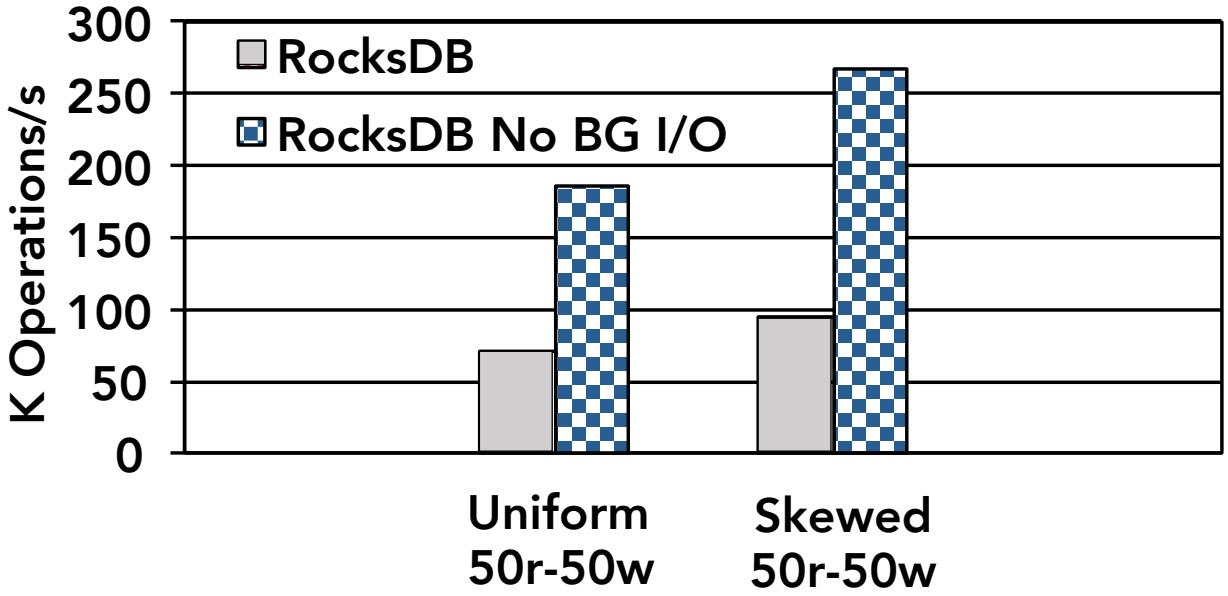

Background I/O Overhead

- **Long & slow bg. ops ⟶ slowdown of user ops.**

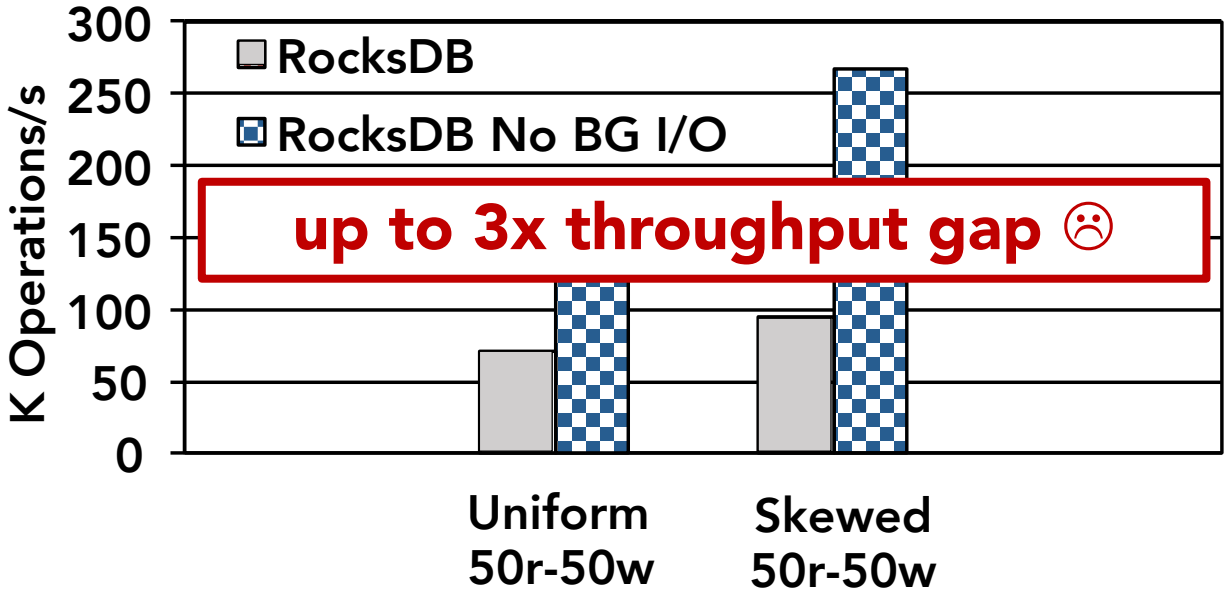

Goal

Decrease background ops overhead
to increase user throughput.

TRIAD

TRIAD

	Workload	Improve WA in
 **TRIAD-MEM**	Skewed workloads	Flushing and Compaction
 **TRIAD-DISK**	In-between	Compaction
 **TRIAD-LOG**	Uniform workloads	Flushing

Three techniques work together and are complementary.

TRIAD

	Workload	Improve WA in
 TRIAD-MEM	Skewed workloads	Flushing and Compaction
 TRIAD-LOG	Uniform workloads	Flushing

TRIAD-MEM

	Workload	Improve WA in
 TRIAD-MEM	Skewed workloads	Flushing and Compaction

Problem: Flushing with Skewed Workloads

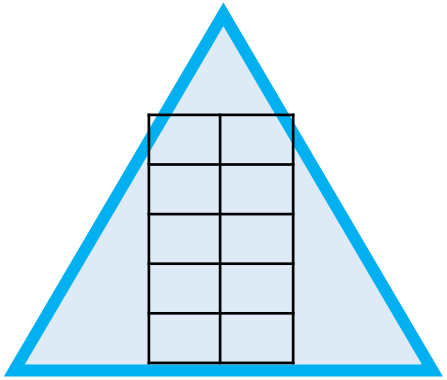

Cm

RAM

Disk

flushing

L0

Commit Log

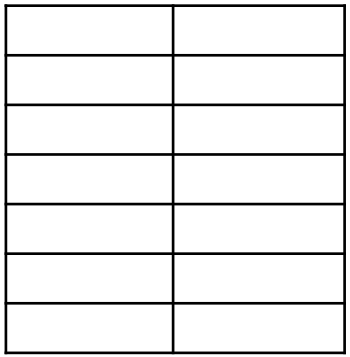

Problem: Flushing with Skewed Workloads

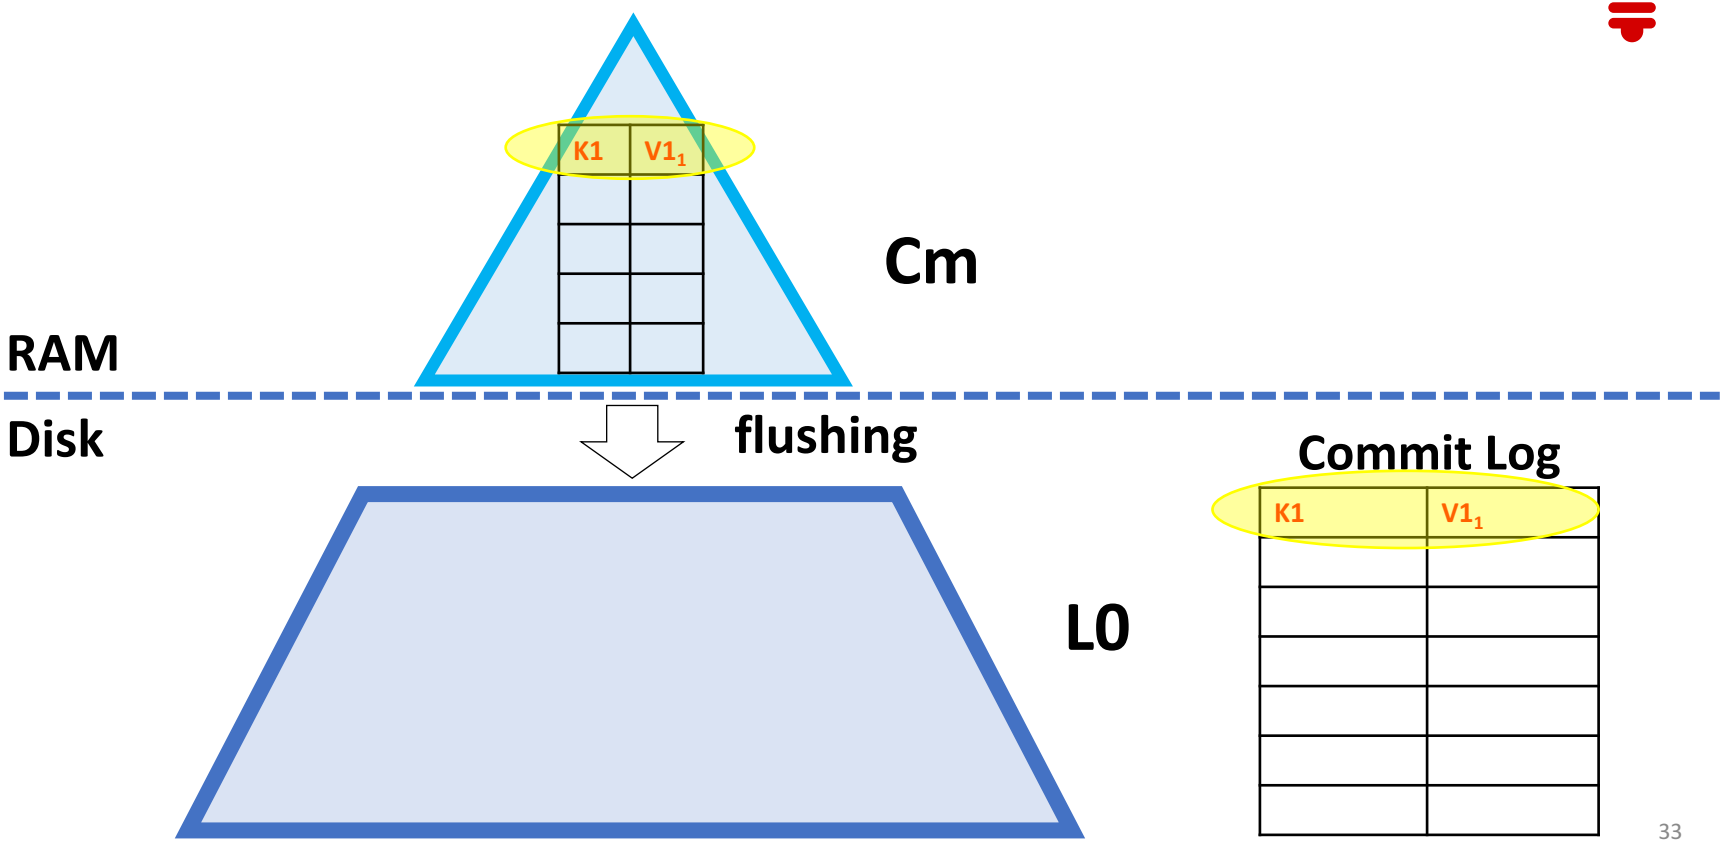

Problem: Flushing with Skewed Workloads

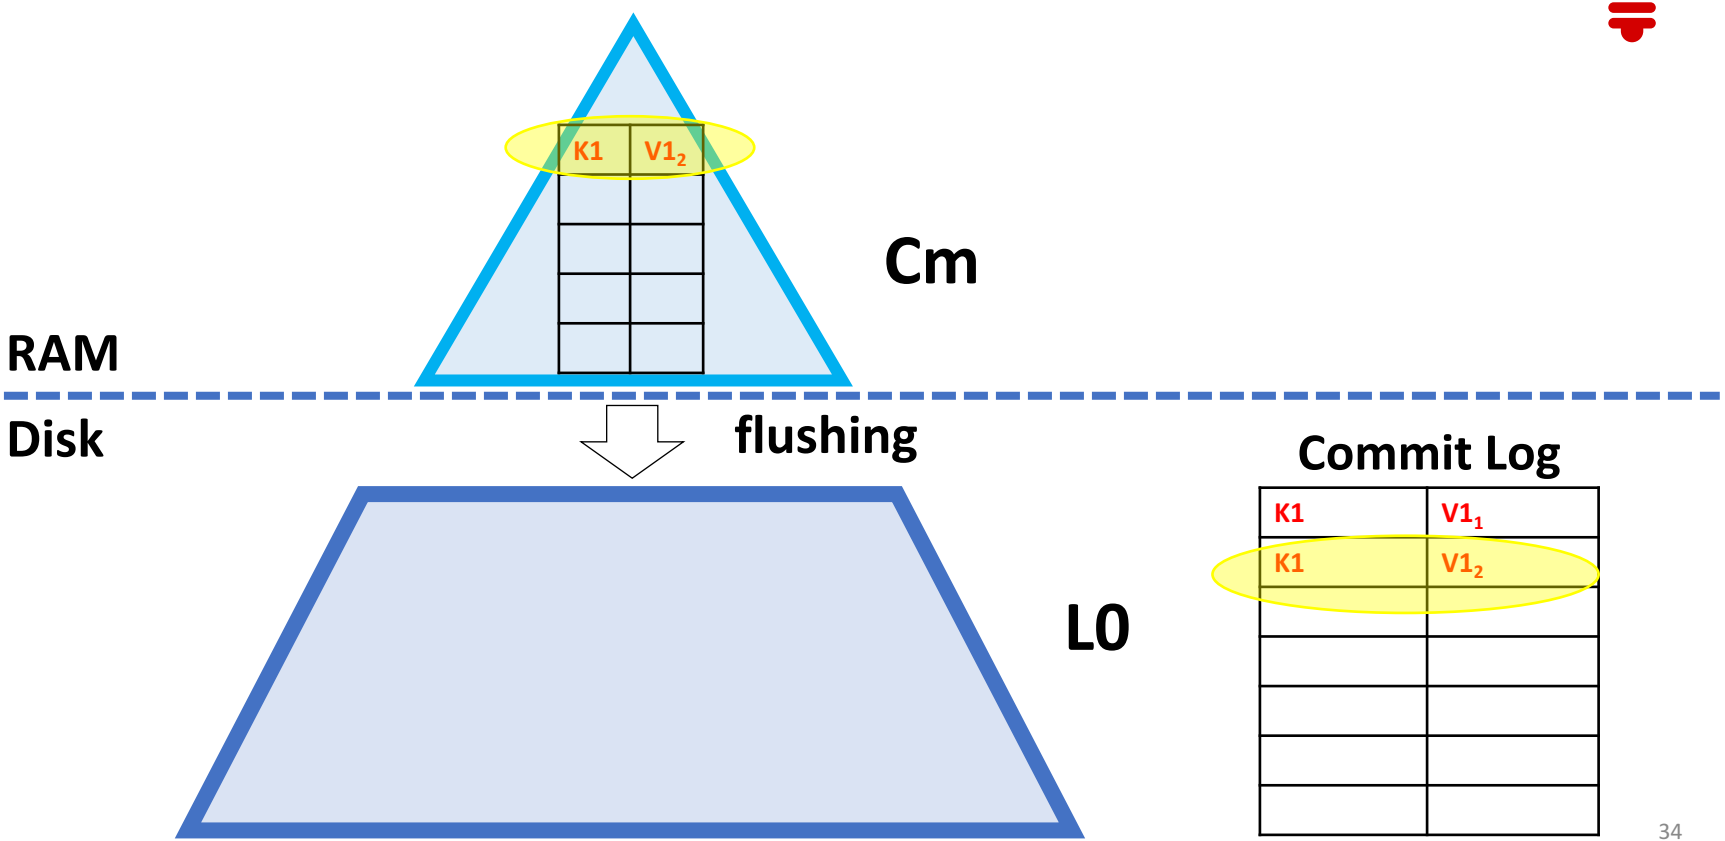

Problem: Flushing with Skewed Workloads

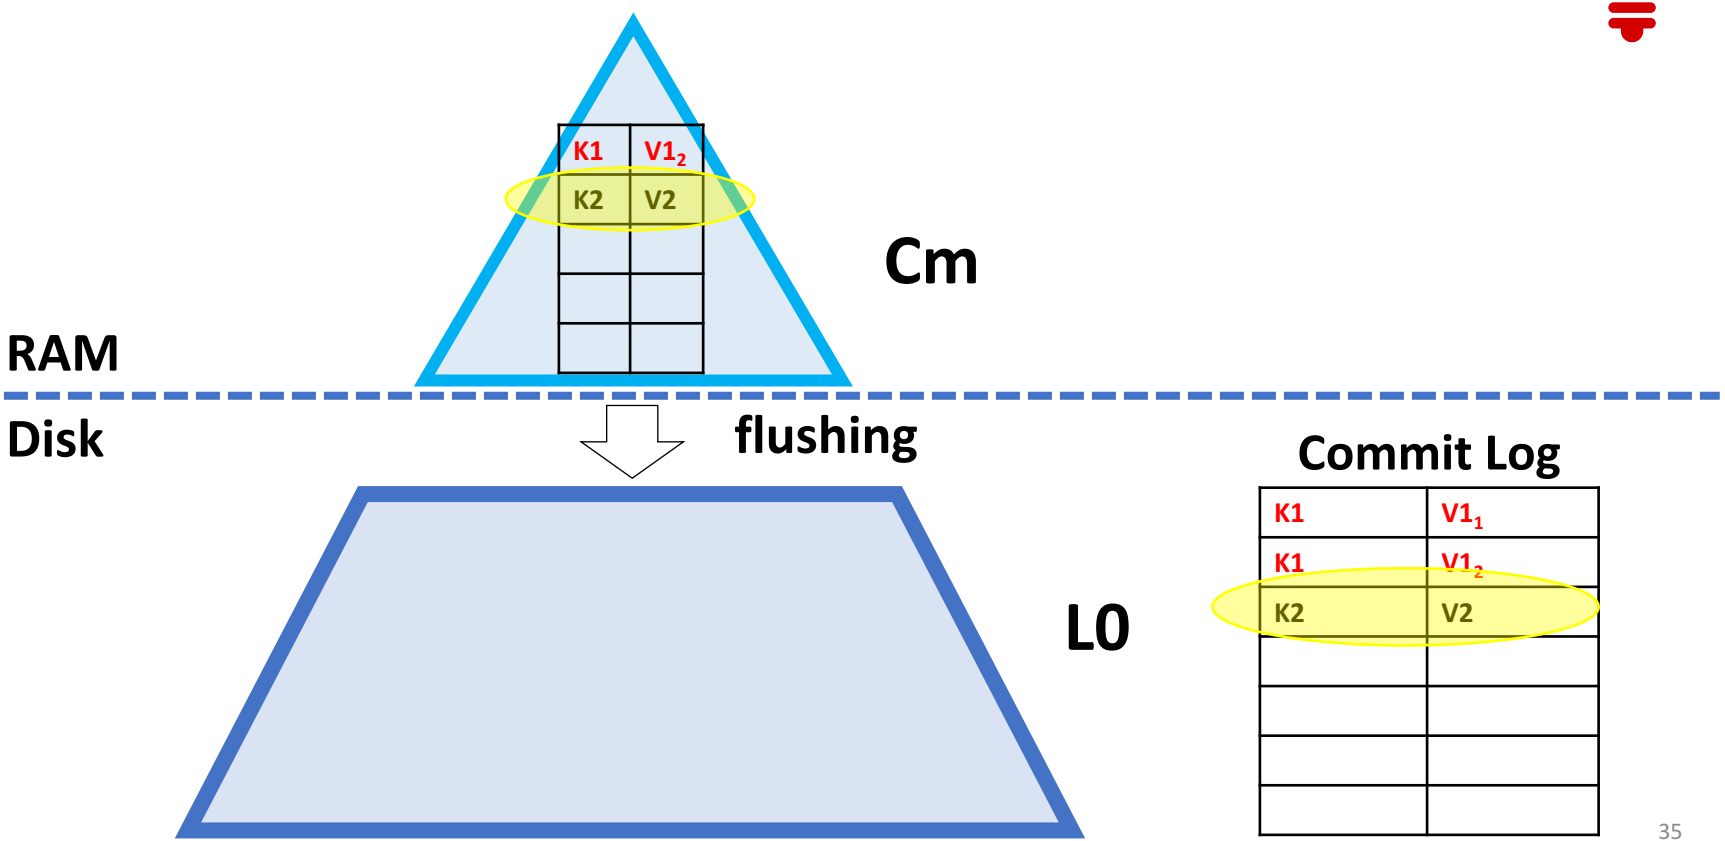

Problem: Flushing with Skewed Workloads

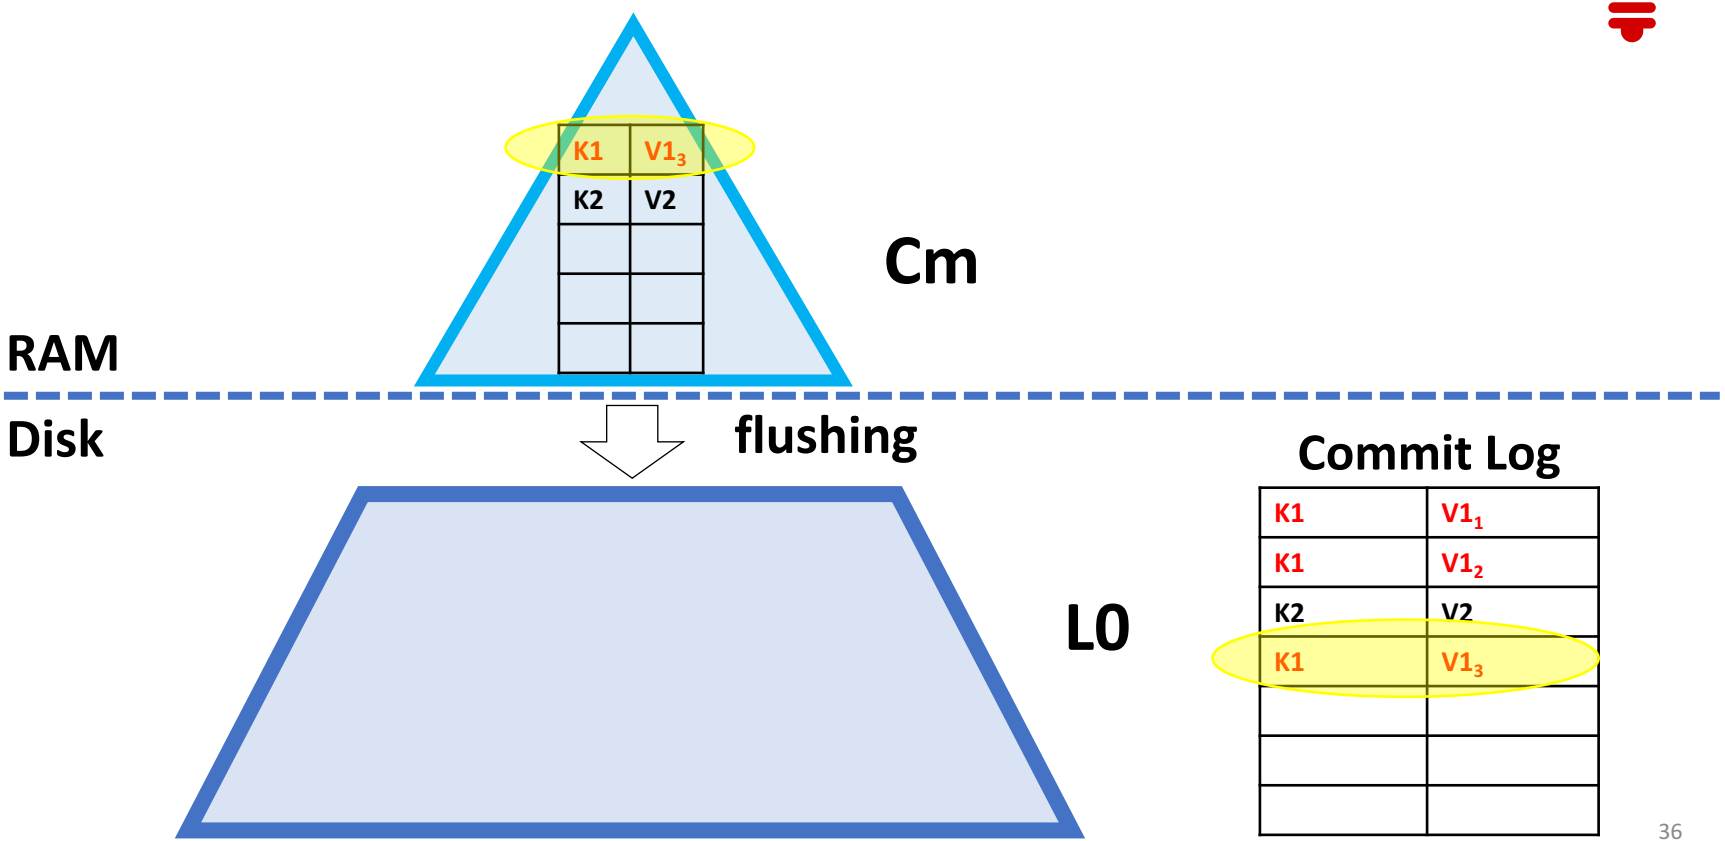

Problem: Flushing with Skewed Workloads

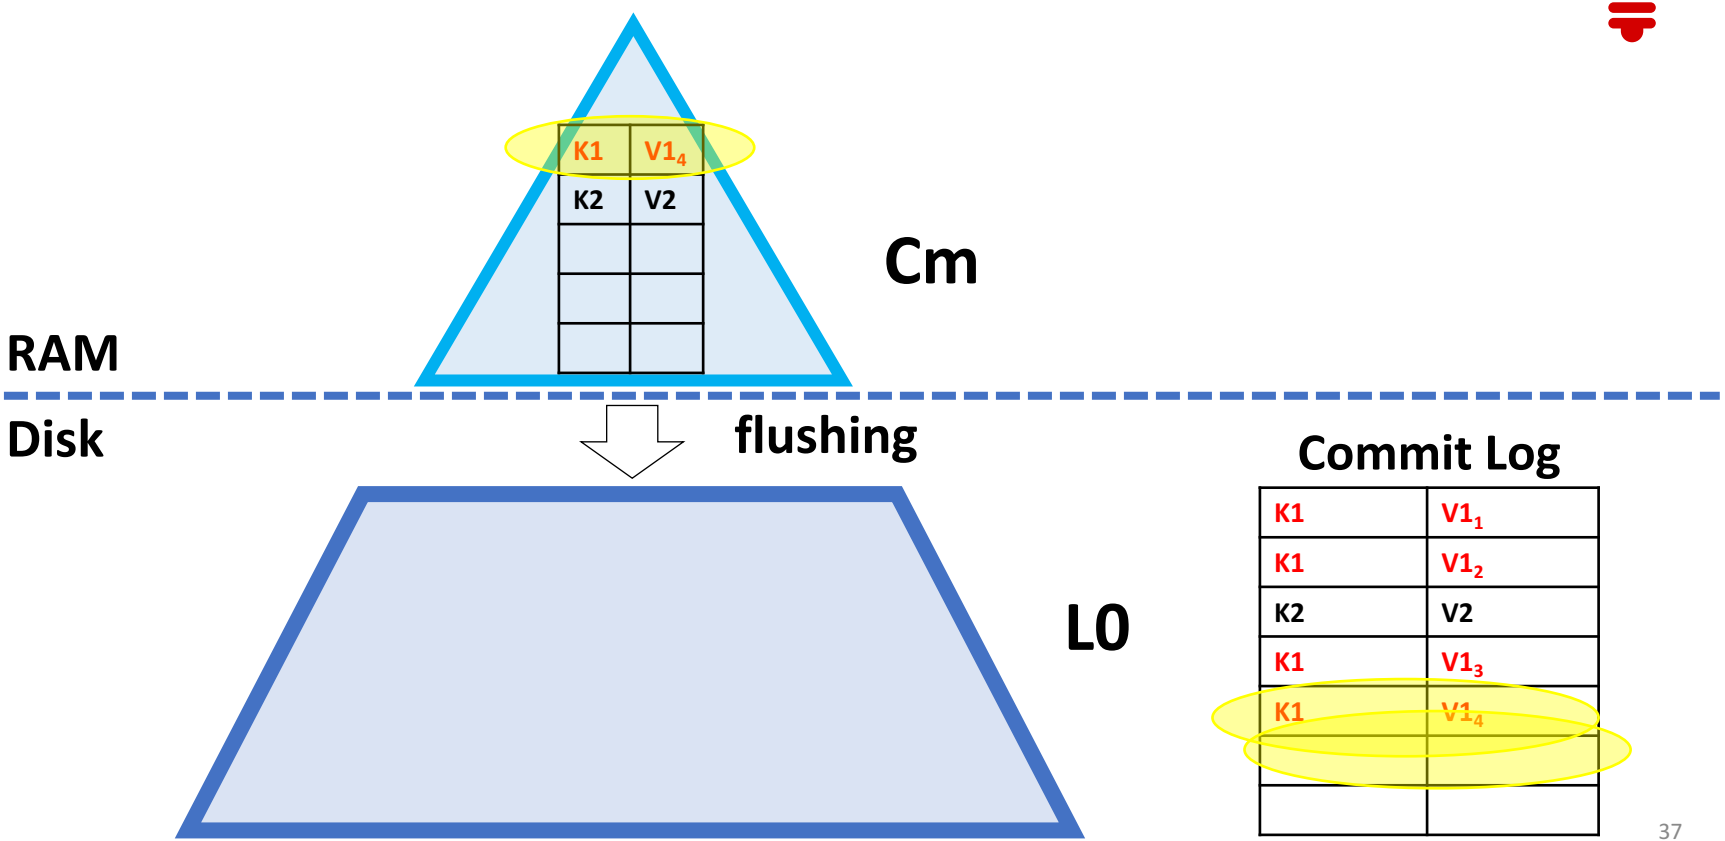

Problem: Flushing with Skewed Workloads

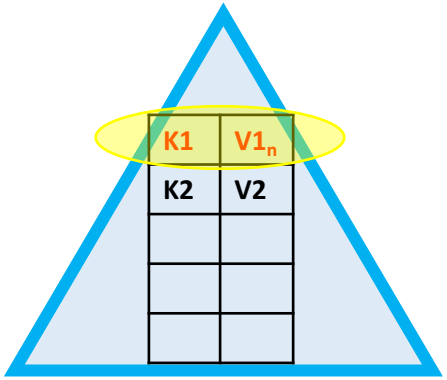

Cm

RAM

Disk

flushing

L0

Commit Log

K1	$V1_1$
K1	$V1_2$
K2	V2
K1	$V1_3$
K1	$V1_4$
…	
K1	$V1_n$

Problem: Flushing with Skewed Workloads

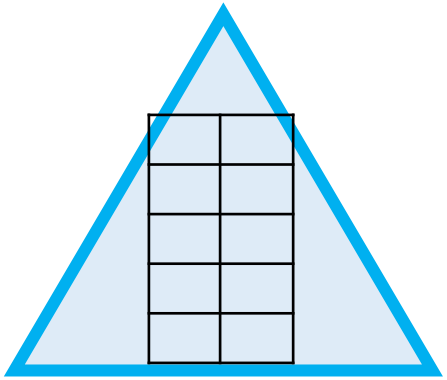

Cm

RAM

Disk

flushing

Commit Log

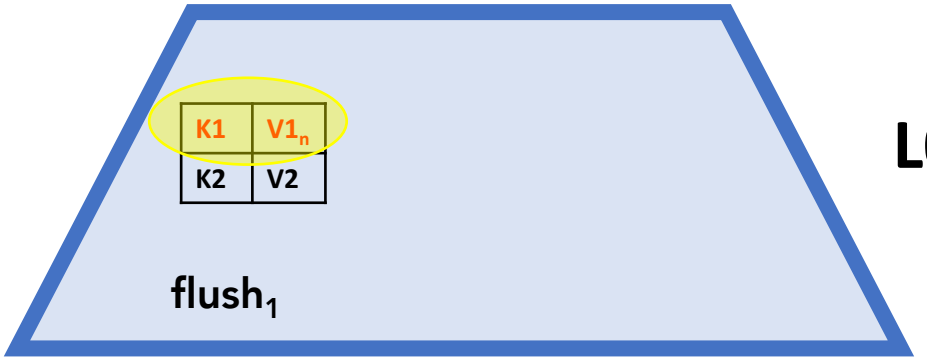

L0

Problem: Flushing with Skewed Workloads

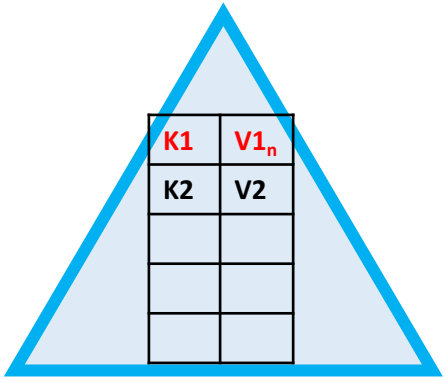

Cm

RAM

Disk

flushing

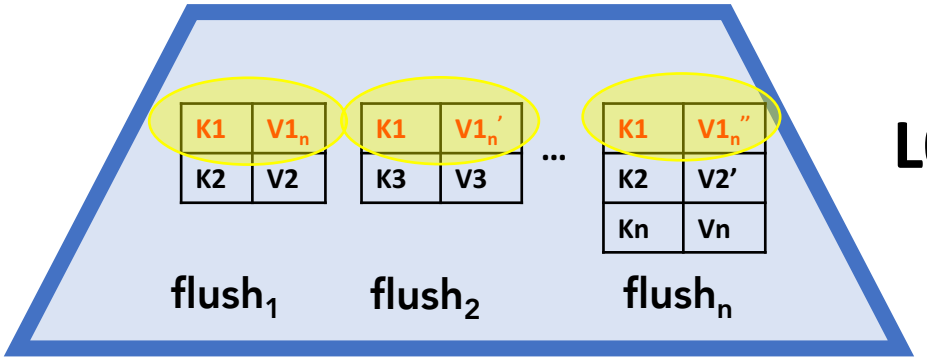

L0

High data skew

- Flush because commit log is full
- Flush mostly empty mem comp

Commit Log

K1	$V1_1$
K2	V2
K1	$V1_2$
K1	$V1_3$
K1	$V1_4$
...	
K1	$V1_n$

Problem: Compaction with Skewed Workloads

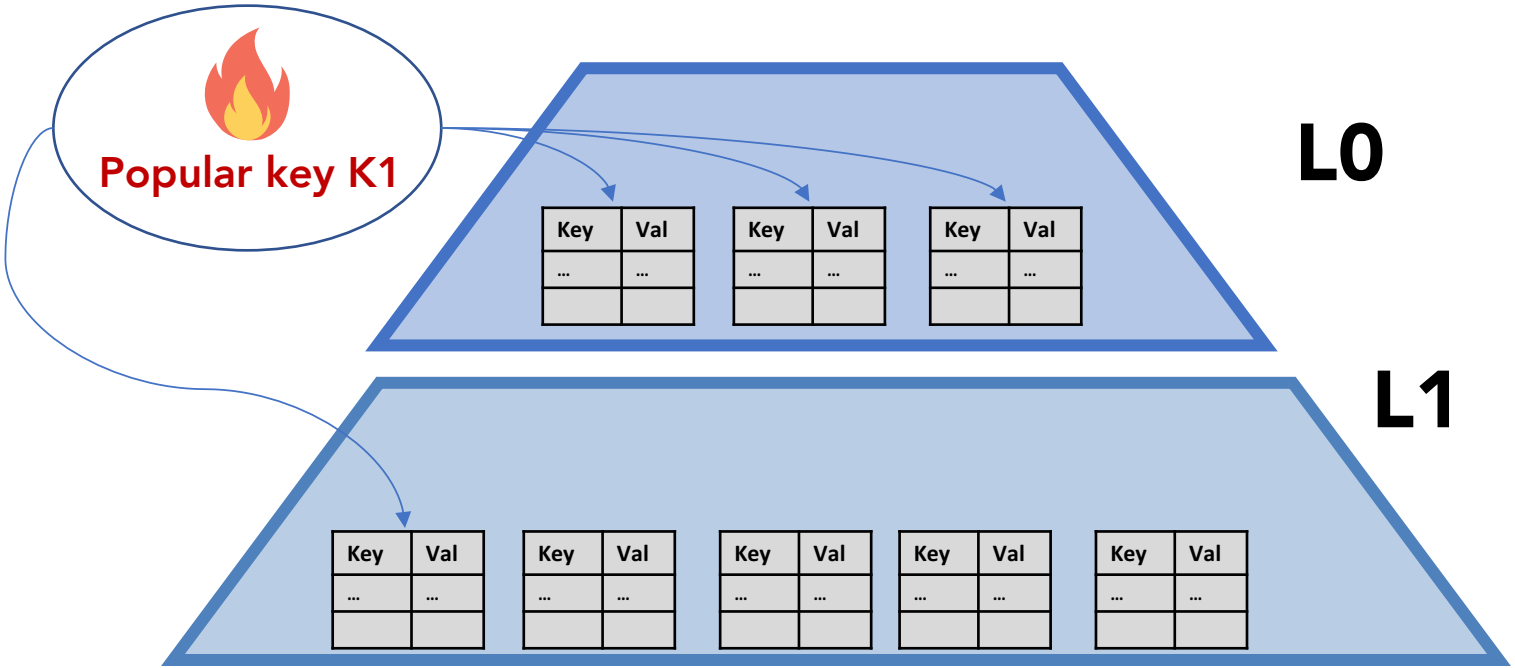

Problem: Compaction with Skewed Workloads

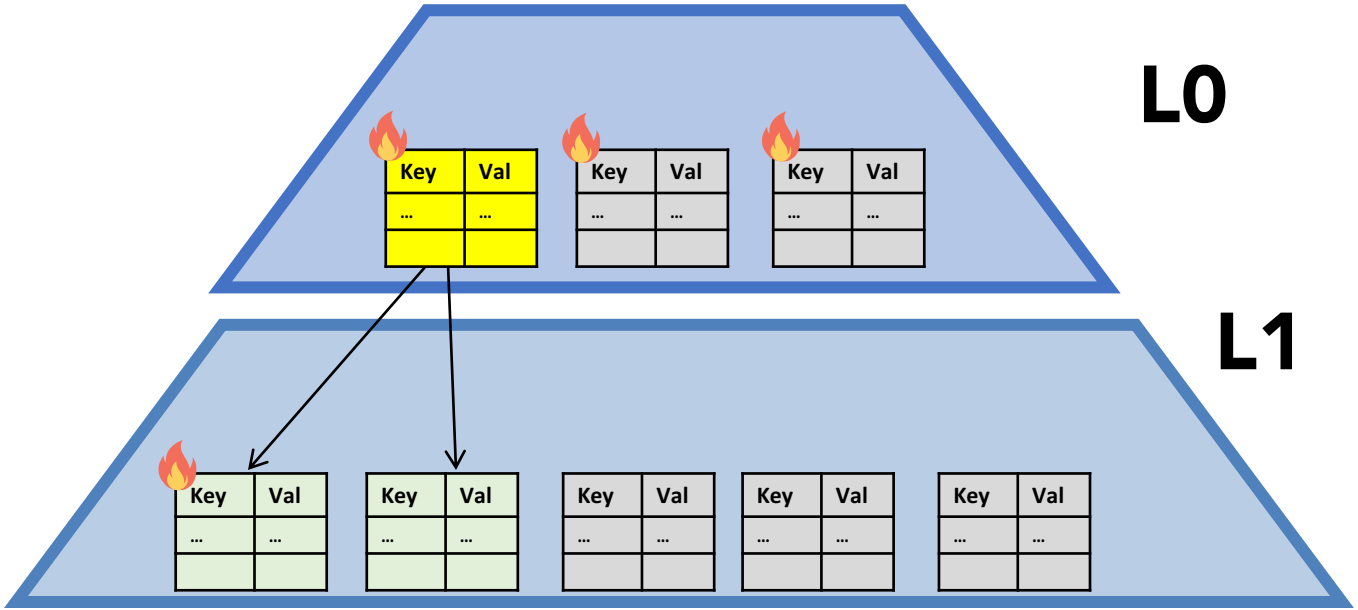

Problem: Compaction with Skewed Workloads

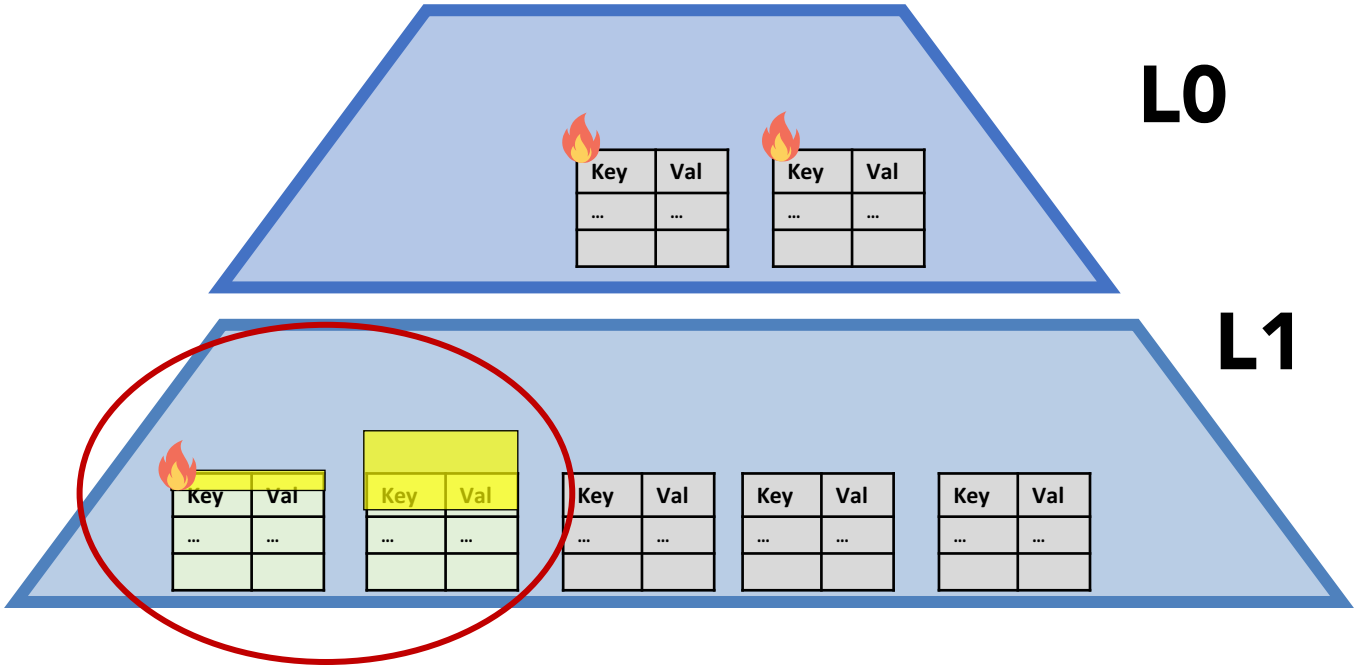

Problem: Compaction with Skewed Workloads

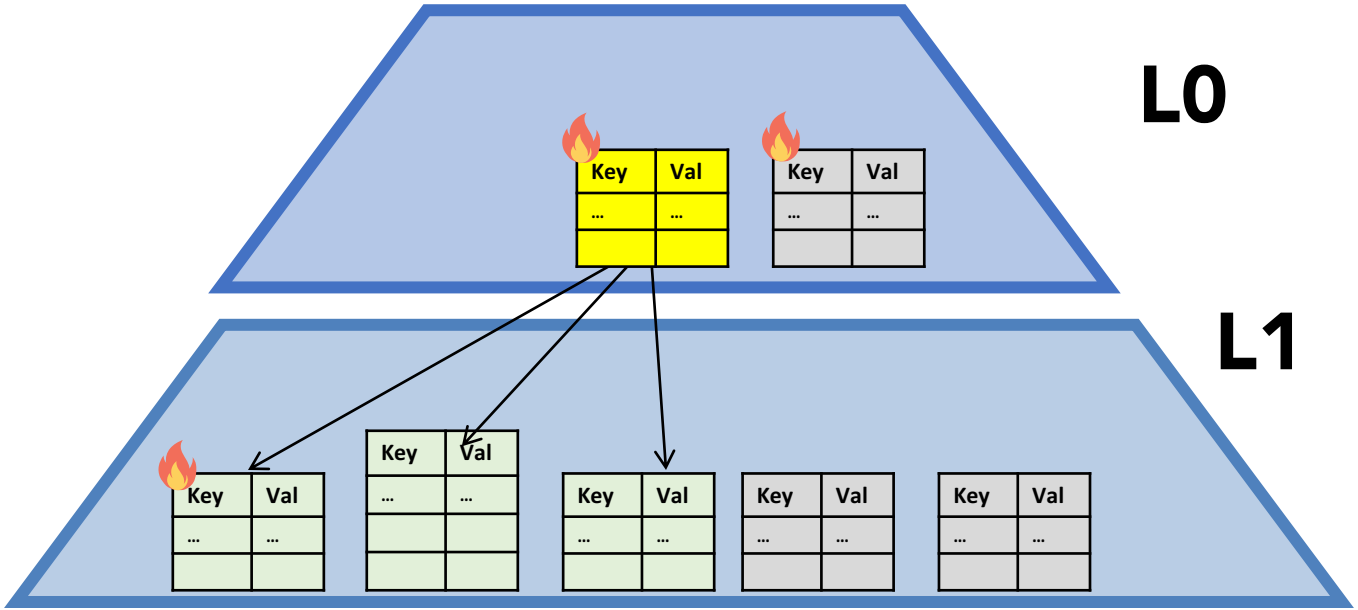

Problem: Compaction with Skewed Workloads

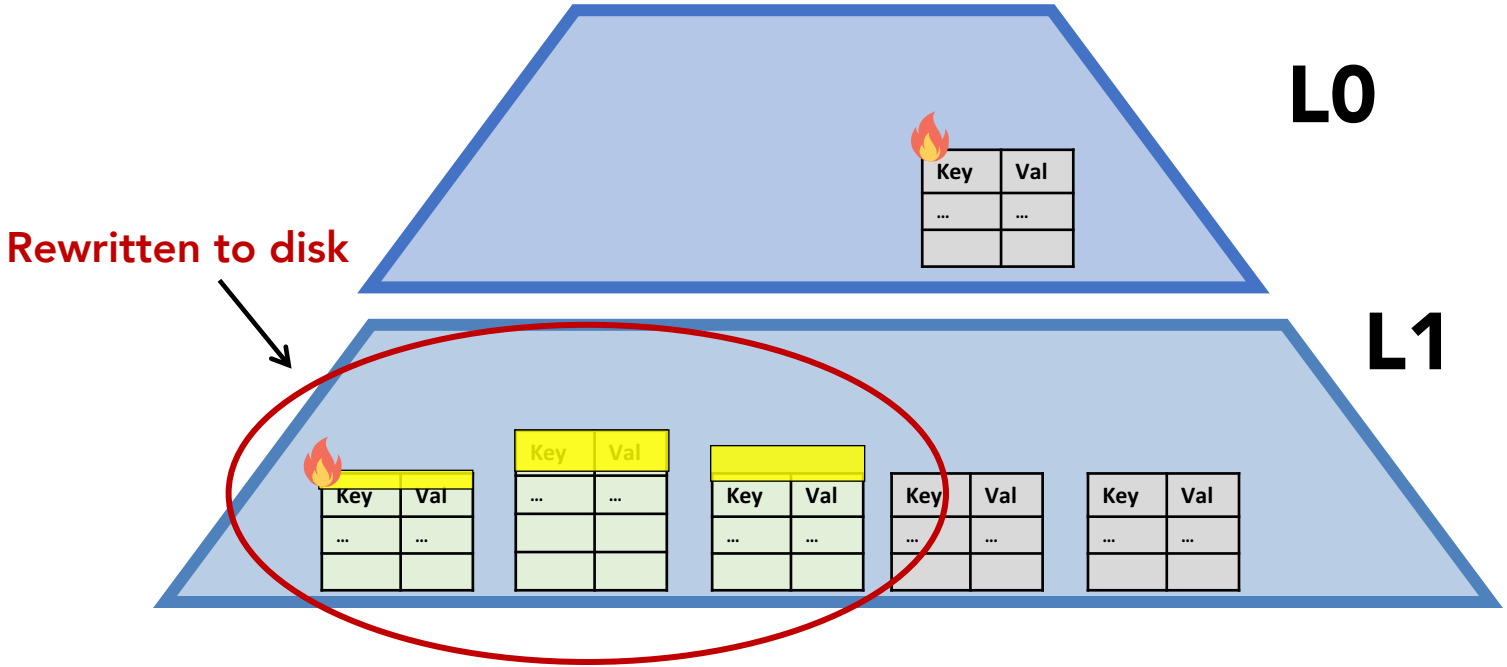

Problem: Compaction with Skewed Workloads

File on L1 **rewritten to disk twice** because of one key ☹

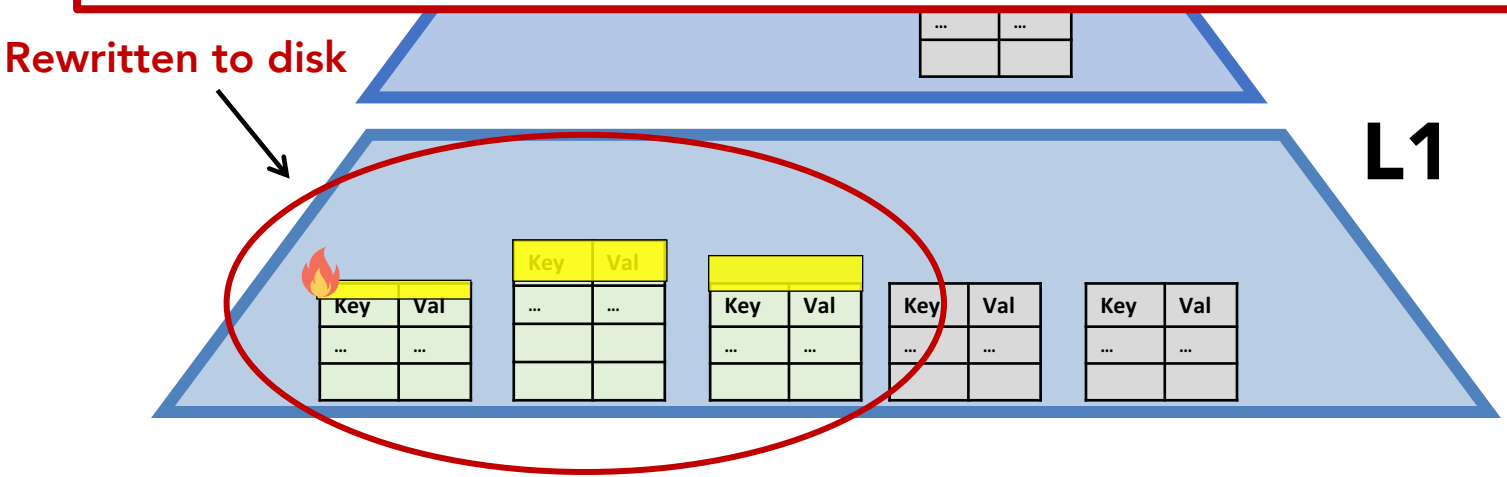

TRIAD-MEM: Hot-cold key separation

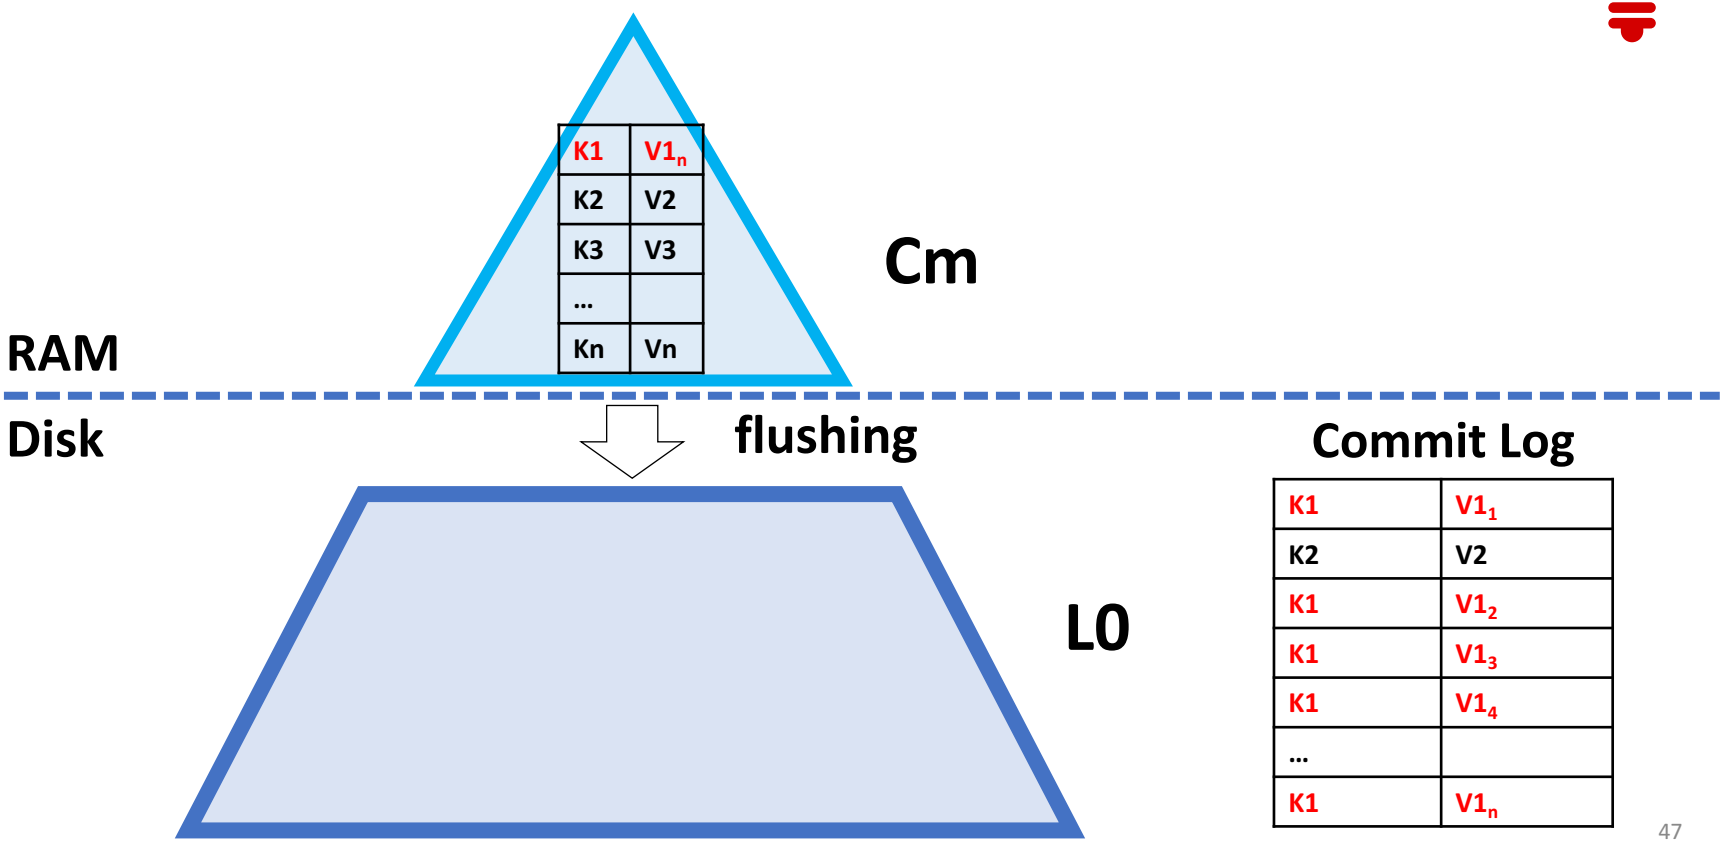

TRIAD-MEM: Hot-cold key separation

Idea:

Keep **hot keys in memory**

Flush only **cold keys**

Keep **hot keys in CL**

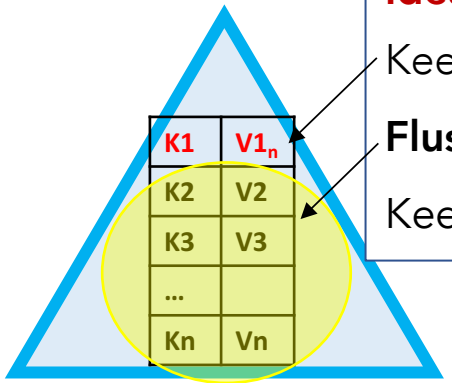

RAM

Disk

flushing

Commit Log

K1	$V1_1$
K2	V2
K1	$V1_2$
K1	$V1_3$
K1	$V1_4$
...	
K1	$V1_n$

TRIAD-MEM: Hot-cold key separation

Idea:

Keep **hot keys in memory**

Flush only **cold keys**

Keep **hot keys in CL**

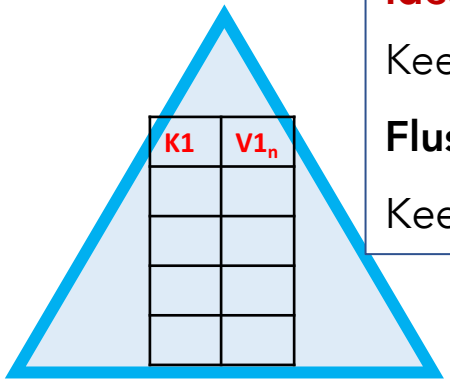

RAM

Disk

flushing

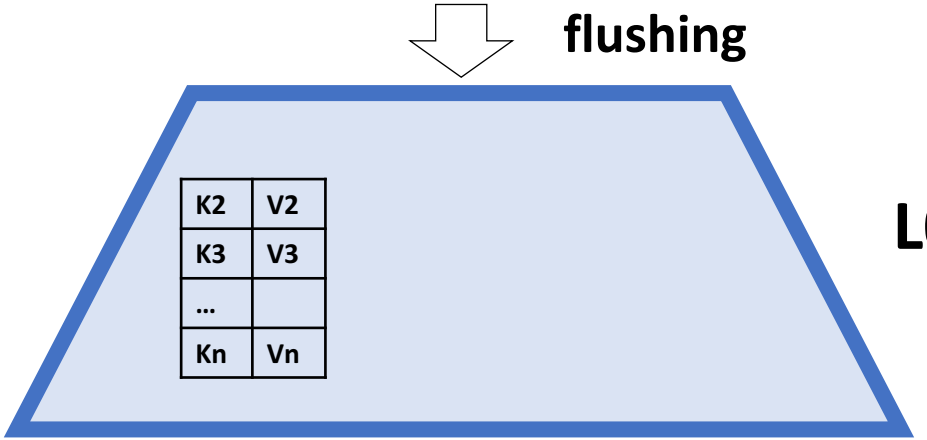

L0

Commit Log

K1	$V1_n$

TRIAD-MEM Summary

- ✓ **Good for skewed workloads.**
- ✓ **Reduce flushing WA: less data written from memory to disk.**
- ✓ **Reduce compaction WA: avoid repeatedly compacting hot keys.**

TRIAD-LOG

	Workload	Improve WA in
 TRIAD-LOG	Uniform workloads	Flushing

Problem: Flushing with Uniform Workloads

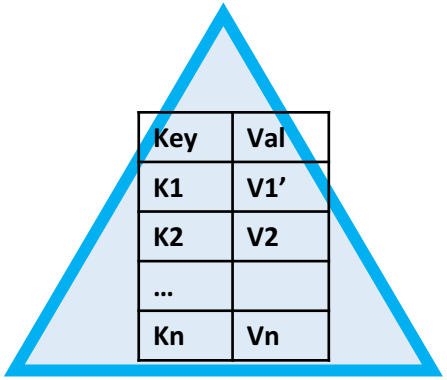

Key	Val
K1	V1’
K2	V2
…	
Kn	Vn

Cm

RAM

Disk

flushing

L0

Commit Log

K1	V1
K2	V2
K1	V1’
K3	V3
K3	V3’
…	
Kn	Vn

Problem: Flushing with Uniform Workloads

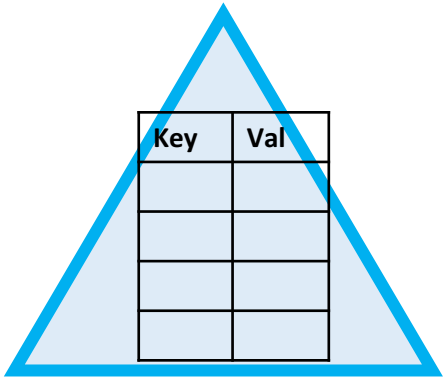

Cm

RAM

Disk

flushing

Key	Val
K1	V1’
K2	V2
…	
Kn	Vn

flush

L0

Commit Log

K1	V1
K2	V2
K1	V1’
K3	V3
K3	V3’
…	
Kn	Vn

Problem: Flushing with Uniform Workloads

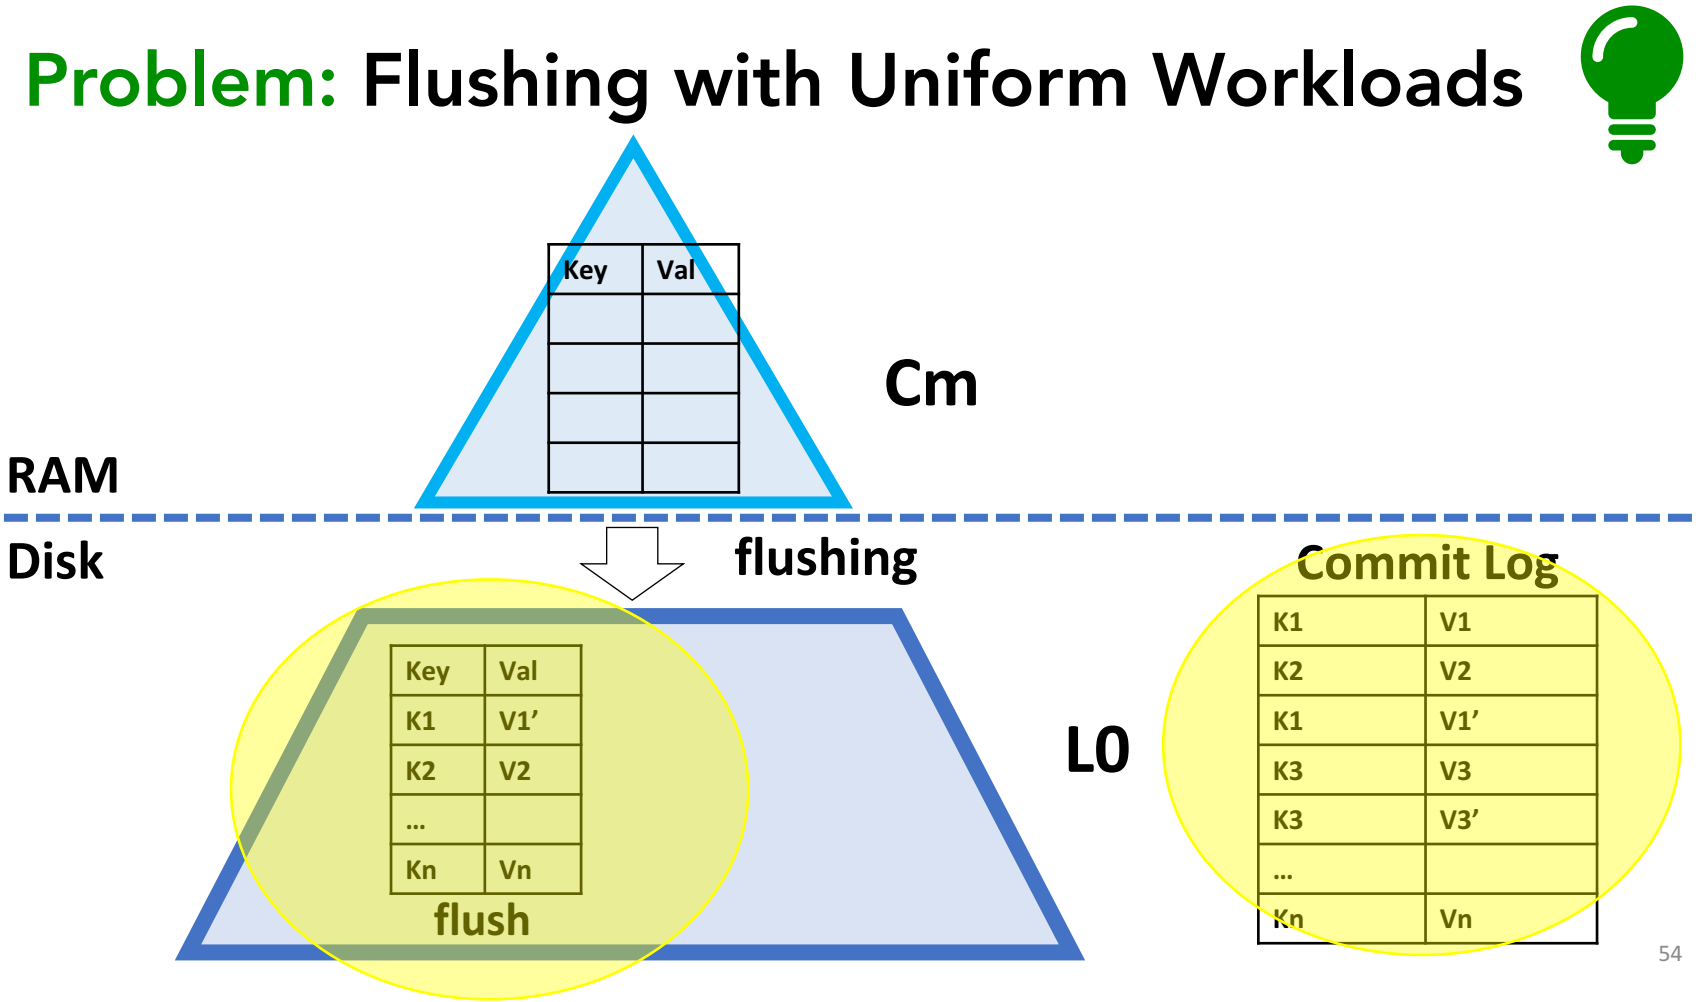

Problem: Flushing with Uniform Workloads

Insight:
Flushed data already written **to commit log.**

Idea:
Use commit logs as SSTables. **Avoid bg I/O due to flushing.**

RAM

Disk

Key	Val
K1	V1’
K2	V2
…	
Kn	Vn

flush

L0

K2	V2
K1	V1’
K3	V3
K3	V3’
…	
Kn	Vn

TRIAD-LOG

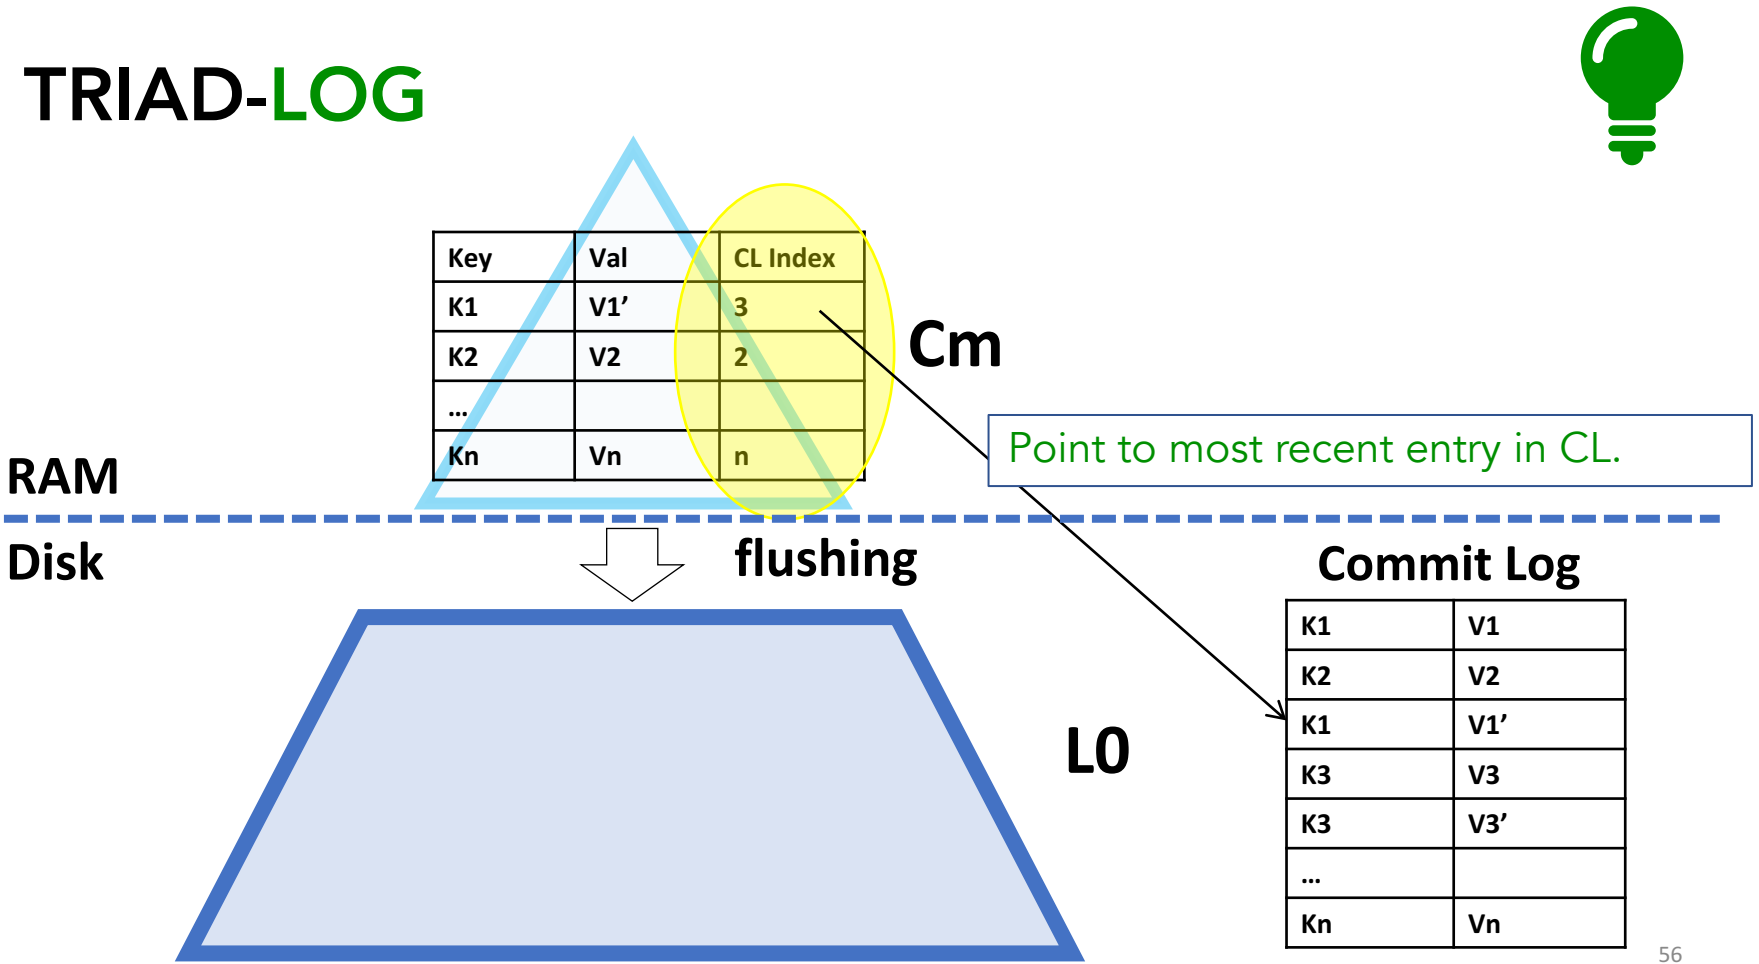
Key
Val
CL Index
K1
V1'
3
K2
V2
2
...
Kn
Vn
n
Cm
RAM
Disk
flushing
L0
Point to most recent entry in CL.
Commit Log
K1
V1
K2
V2
K1
V1'
K3
V3
K3
V3'
...
Kn
Vn

TRIAD-LOG

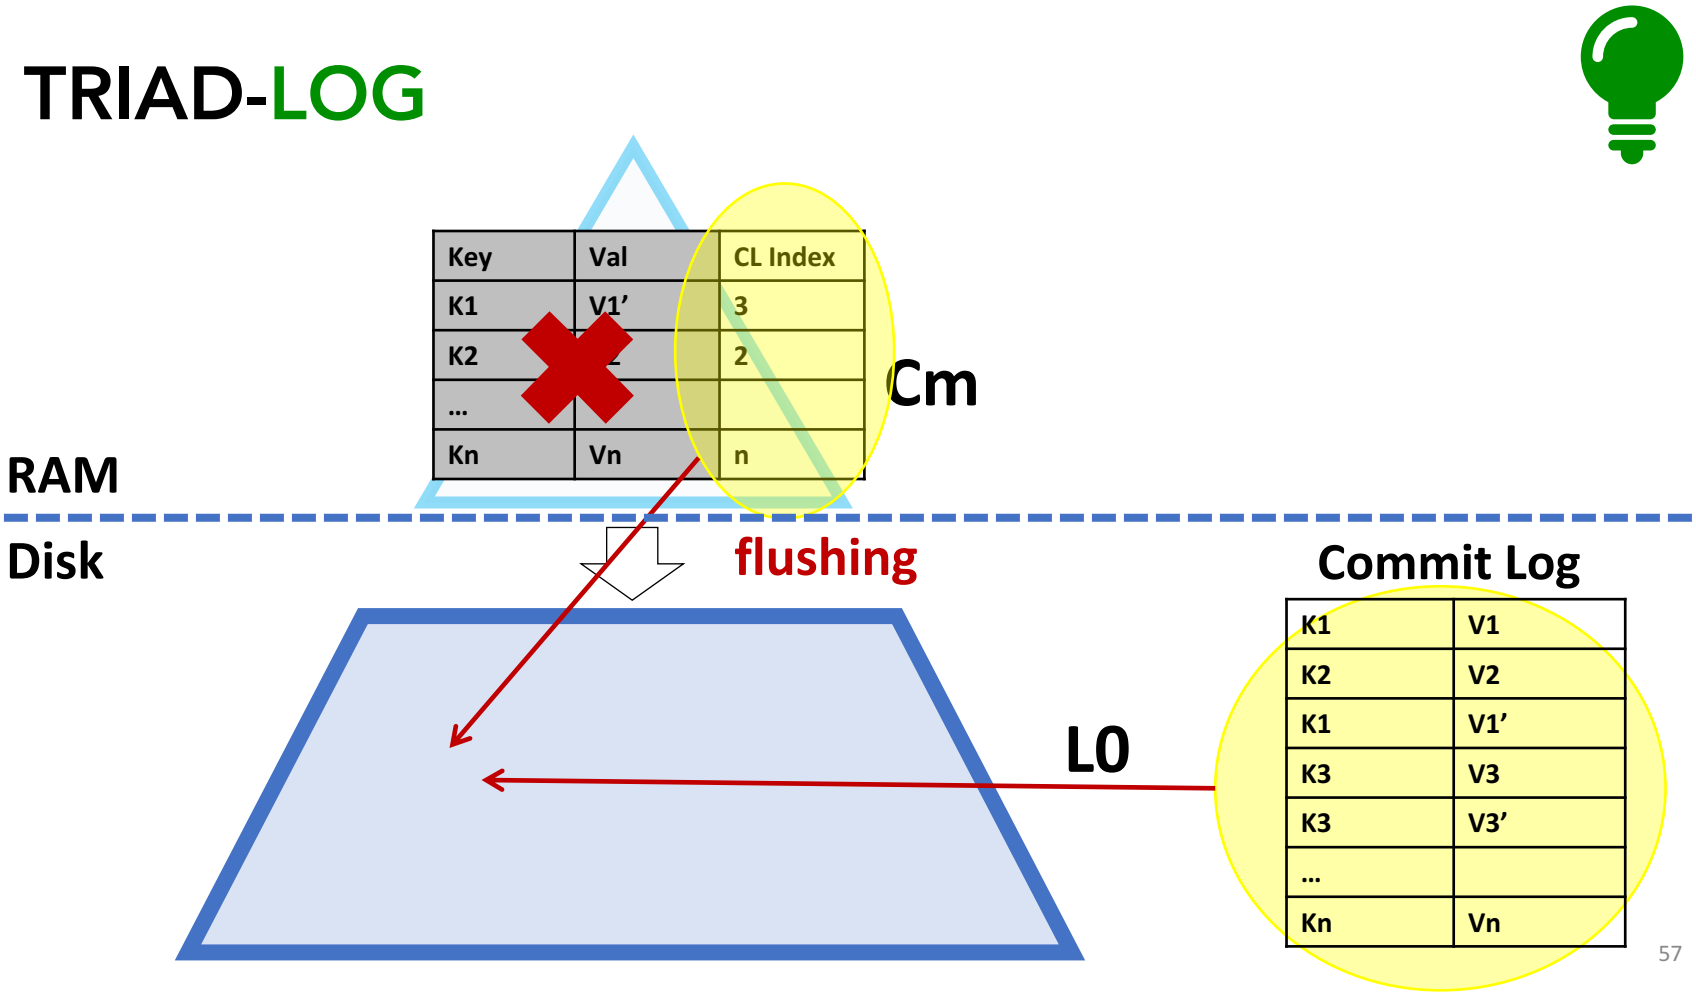
Key
Val
CL Index
K1
V1'
3
K2
2
...
Kn
Vn
n
Cm
RAM
Disk
flushing
Commit Log
K1
V1
K2
V2
K1
V1'
K3
V3
K3
V3'
...
Kn
Vn
L0

TRIAD-LOG

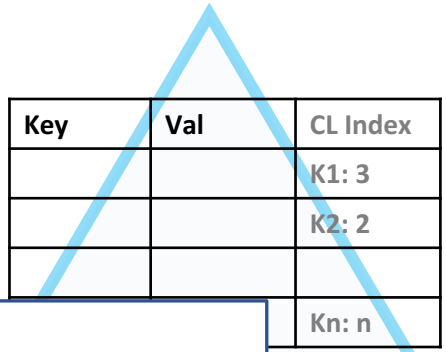

Cm

CL Index
K1: 3
K2: 2
Kn: n

Keep index in memory for further reads.

CL-SSTable
Only flush CL Index from memory and couple it with the current Commit Log.

flushing

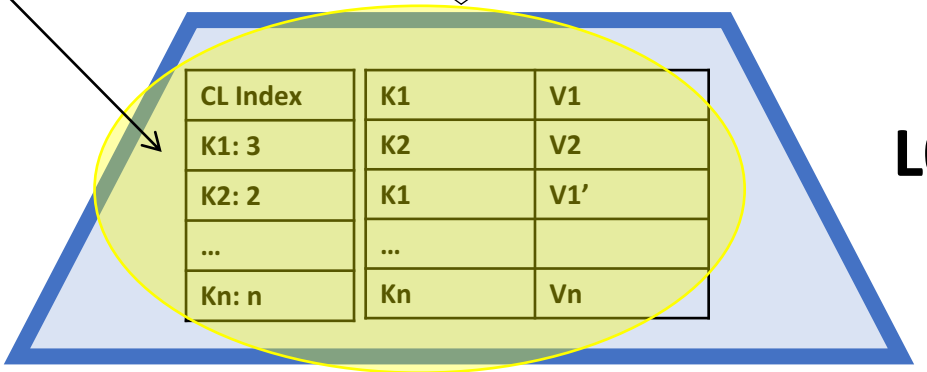

L0

Commit Log

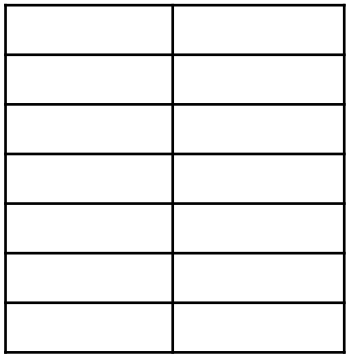

TRIAD-LOG Summary

- ✓ **Good for uniform workloads.**
- ✓ **Reuse Commit Log as L0 SST.**
- ✓ **No more flushing of mem component to disk.**

TRIAD Summary

TRIAD-MEM, TRIAD-DISK, TRIAD-LOG:

- **Complementary**, targeting different wklds.
- Working **simultaneously**.
- **Transparent** to the workloads; no a priori knowledge needed.

Evaluation

Evaluation

- **Compare TRIAD with RocksDB**
- **Workloads:** Production, Synthetic
- **Metrics:** Throughput, Write Amplification (WA)
- **Code:** https://github.com/epfl-labos/TRIAD

Write Amplification (WA)

$$WA = \frac{\text{total data written to storage}}{\text{data written by app}}$$

Production Workloads: Throughput

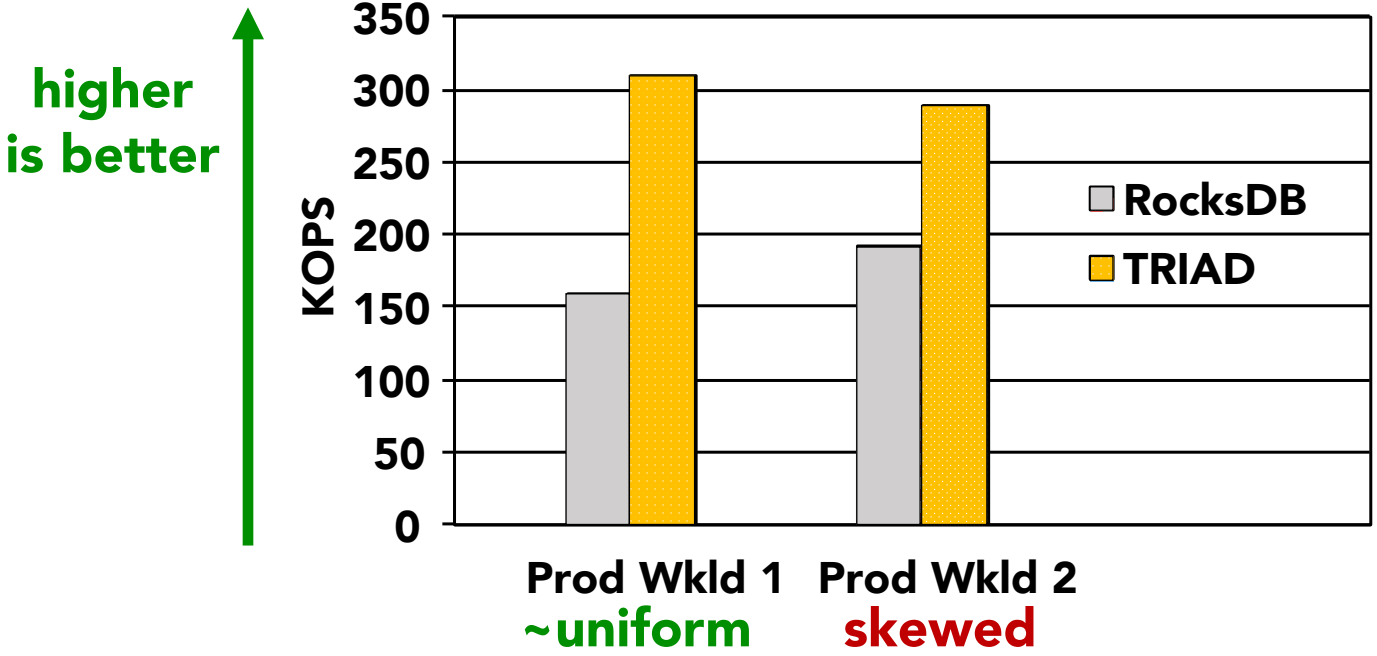

Production Workloads: Throughput

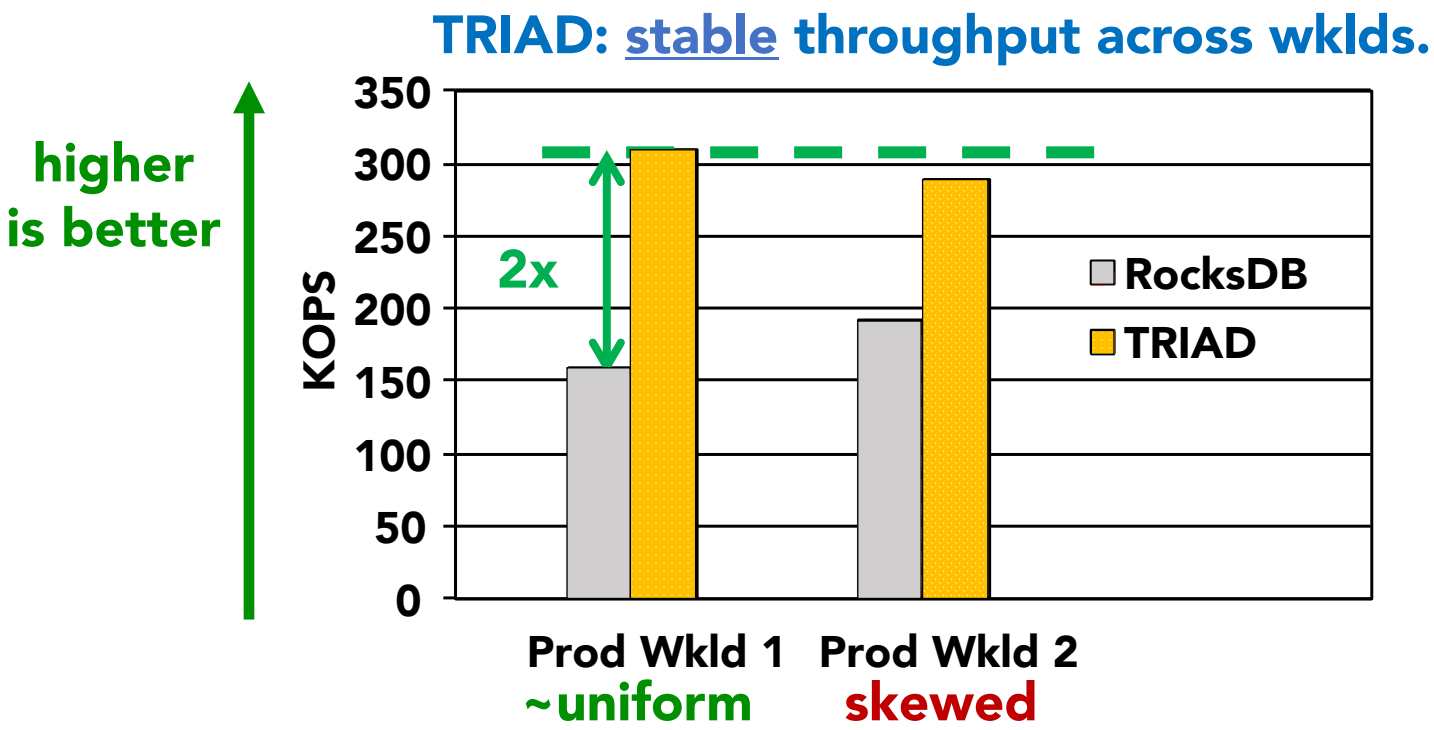

Production Workloads: Write Amplification

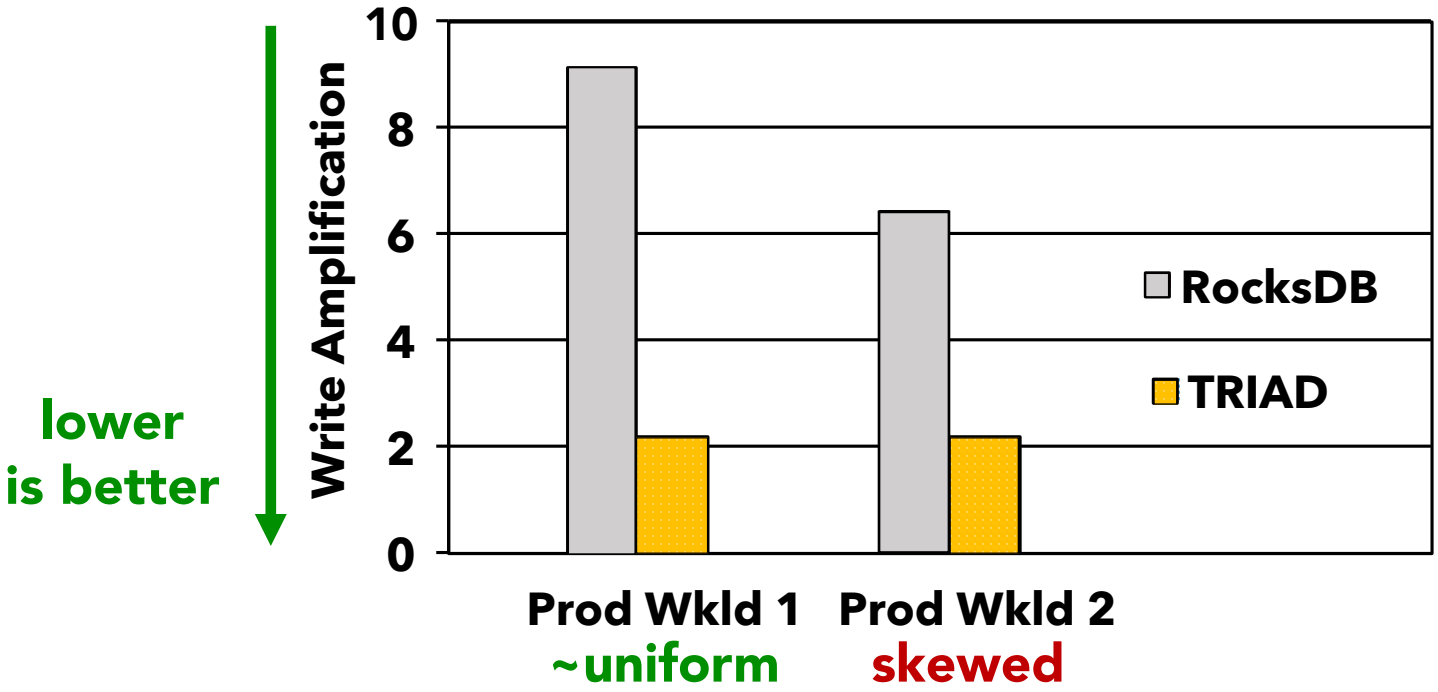

Production Workloads: Write Amplification

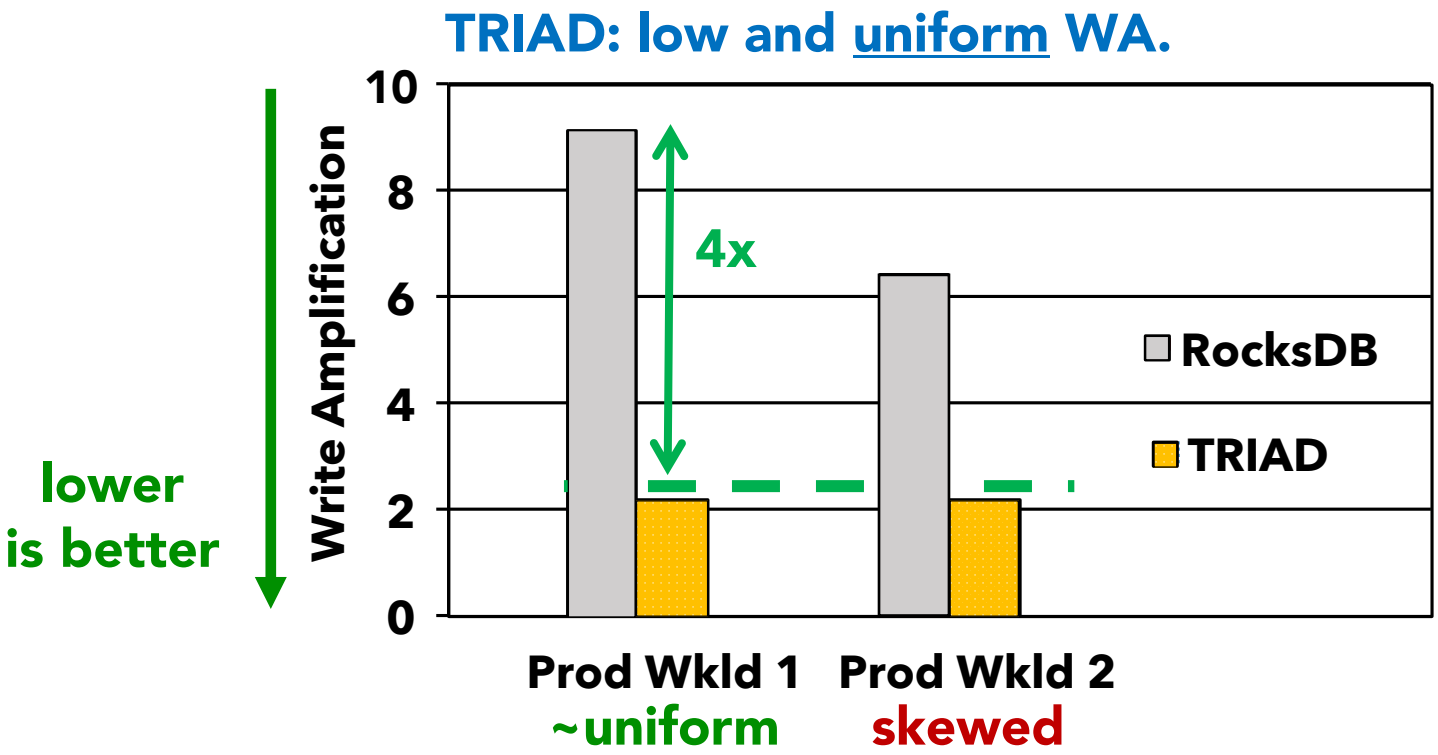

TRIAD: Throughput Breakdown Synthetic Workloads

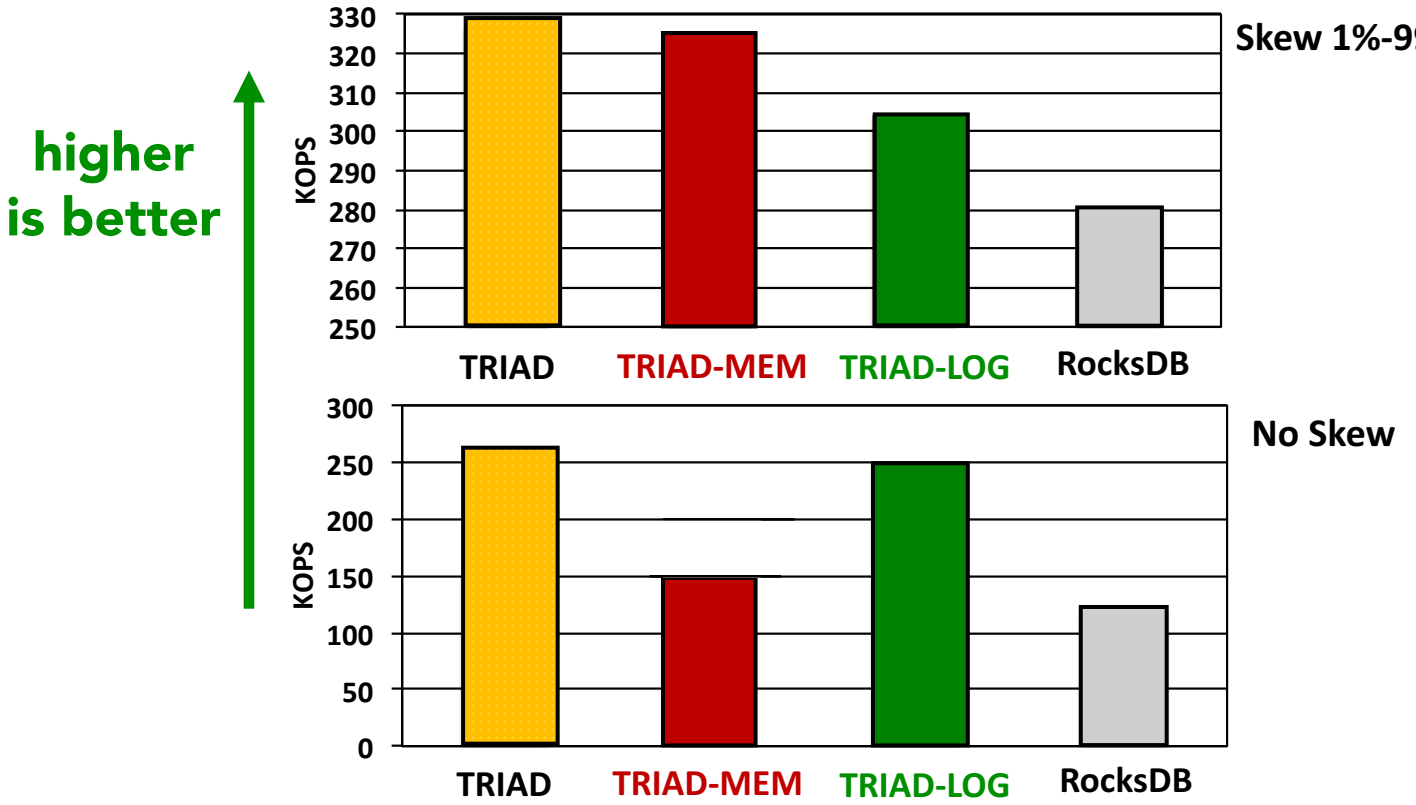

TRIAD: Throughput Breakdown Synthetic Workloads

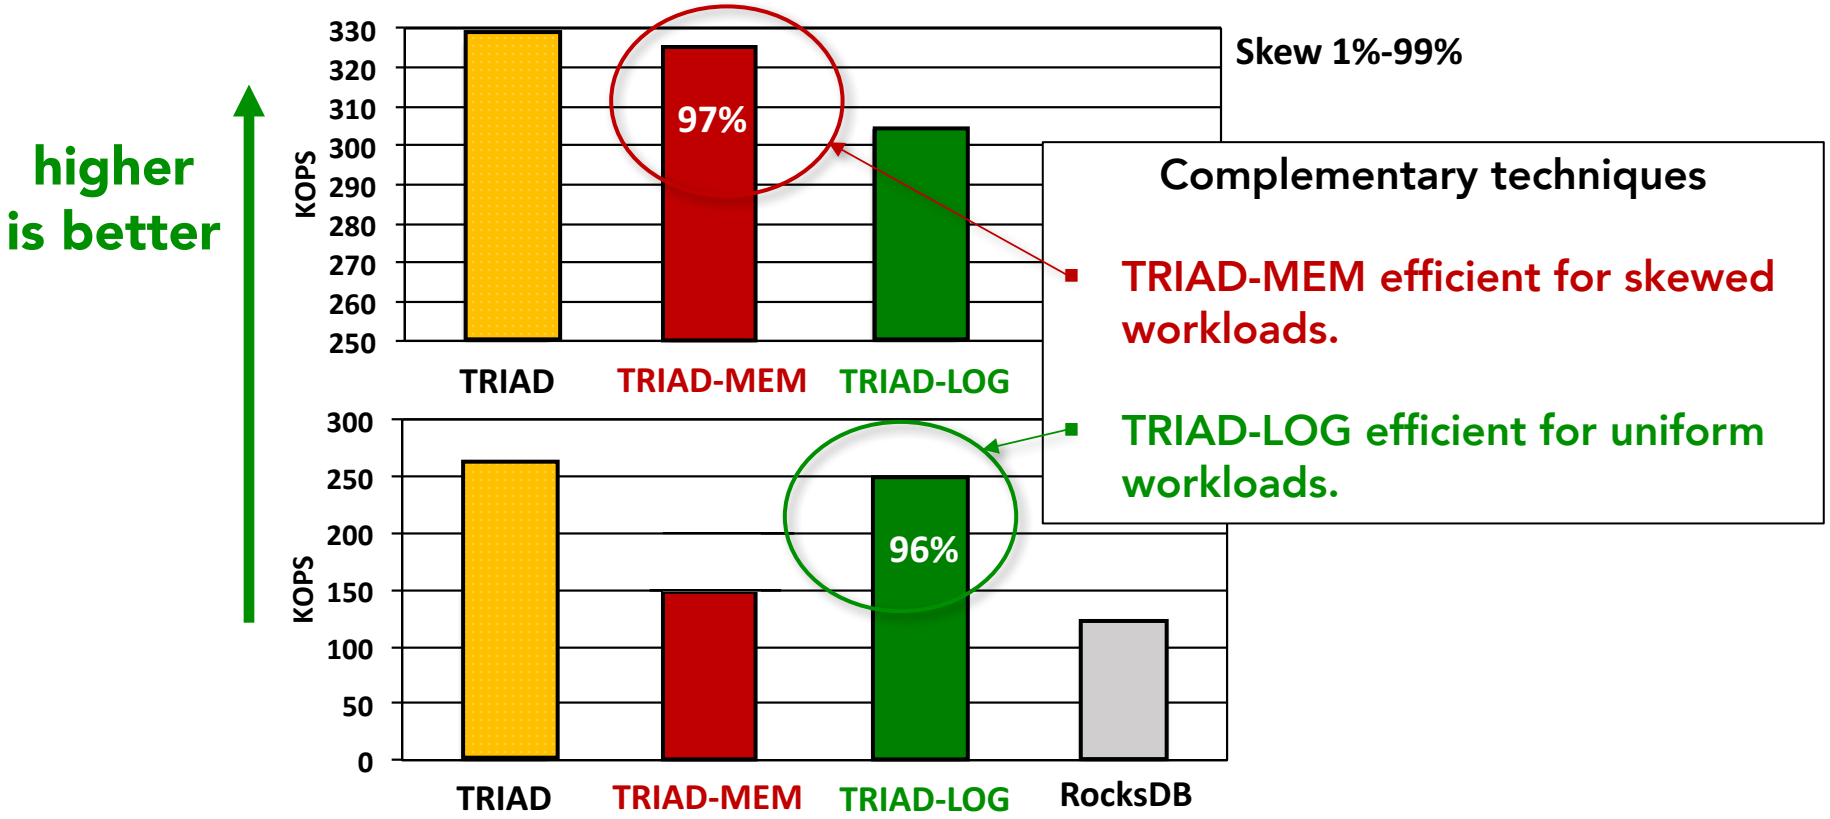

More in Our Paper

- More production workloads
- More synthetic workloads
- Detailed breakdown of TRIAD techniques
- TRIAD-DISK

Related work

- **LevelDB**: first LSM-based KV store. No attempts to reduce WA.

Related work

- **LevelDB**: first LSM-based KV store. No attempts to reduce WA.

- A number of systems attempt to **reduce WA.**
  - **LSMs:** RocksDB, *bLSM* (SIGMOD/PODS '12), *VT-tree* (FAST '13), *HyperLevelDB*, *LSM-trie* (USENIX ATC '15), Cassandra, WiscKey (FAST '16).

Related work

- **LevelDB**: first LSM-based KV store. No attempts to reduce WA.

- A number of systems attempt to **reduce WA.**
  - **LSMs:** RocksDB, *bLSM* (SIGMOD/PODS '12), *VT-tree* (FAST '13), *HyperLevelDB*, *LSM-trie* (USENIX ATC '15), Cassandra, WiscKey (FAST '16).
  - **B-epsilon trees:** Tucana (USENIX ATC '16), BetrFS (FAST '15, '16)

Related work

- **LevelDB**: first LSM-based KV store. No attempts to reduce WA.

- A number of systems attempt to **reduce WA.**
  - **LSMs:** RocksDB, *bLSM* (SIGMOD/PODS '12), *VT-tree* (FAST '13), *HyperLevelDB*, *LSM-trie* (USENIX ATC '15), Cassandra, WiscKey (FAST '16).
  - **B-epsilon trees:** Tucana (USENIX ATC '16), BetrFS (FAST '15, '16)
  - But no hot/cold key separation, not use commit log as a pseudo-SST.

Related work

- **LevelDB**: first LSM-based KV store. No attempts to reduce WA.
- A number of systems attempt to **reduce WA.**
  - **LSMs:** RocksDB, *bLSM* (SIGMOD/PODS '12), *VT-tree* (FAST '13), *HyperLevelDB*, *LSM-trie* (USENIX ATC '15), Cassandra, WiscKey (FAST '16).
  - **B-epsilon trees:** Tucana (USENIX ATC '16), BetrFS (FAST '15, '16)
  - But no hot/cold key separation, not use commit log as a pseudo-SST.
- Hot/cold key separation idea used in different context.
  - **FLASH storage systems:** dual-pool algorithm (SAC '07), Application-Managed Flash (FAST '16)

Take-home Messages

✓ **TRIAD: LSM key-value store with high throughput and low WA.**

Take-home Messages

✓ **TRIAD: LSM key-value store with high throughput and low WA.**

✓ **Complementary techniques transparent to workload types.**

Take-home Messages

✓ **TRIAD: LSM key-value store with high throughput and low WA.**

✓ **Complementary techniques transparent to workload types.**

✓ **Impact of LSM I/O on user throughput reduced.**